

SC3260 / SC5260

Fault-tolerant Techniques for HPC

Lecture by: Ana Gainaru

Table of content

Slides heavily based on the SC tutorial given by George Bosilca

<https://fault-tolerance.org/downloads/sc14tutorial.pdf>

- ▶ Failures and faults in HPC
- ▶ Checkpointing
 - ▶ Theory
 - ▶ System and application level solutions currently used
- ▶ Fault tolerant MPI
- ▶ Designing a resilient application
- ▶ Silent errors



VANDERBILT
UNIVERSITY

Failures and Faults in HPC

Top ranked supercomputers in the US (June 2017)

Rank	Name	Laboratory	Technology	Cores	PFlops/s	MTBF
4	Titan	ORNL	Cray XK7	560,640	17.59	≈ 1 day
5	Sequoia	LLNL	BG/Q	1,572,864	17.17	≈ 1 day
6	Cori	LBNL	Cray XC40	622,336	14.01	≈ 1 day
9	Mira	ANL	BG/Q	786,432	8.59	≈ 1 day

Fail-stop errors Node failure, resource crashes

Silent errors or silent data corruptions (SDCs) Double bit flips, soft faults



VANDERBILT
UNIVERSITY

Failures and Faults in HPC

Top ranked supercomputers in the US (June 2017)

Rank	Name	Laboratory	Technology	Cores	PFlops/s	MTBF
4	Titan	ORNL	Cray XK7	560,640	17.59	≈ 1 day
5	Sequoia	LLNL	BG/Q	1,572,864	17.17	≈ 1 day
6	Cori	LBNL	Cray XC40	622,336	14.01	≈ 1 day
9	Mira	ANL	BG/Q	786,432	8.59	≈ 1 day

Fail-stop errors Node failure, resource crashes

Silent errors or silent data corruptions (SDCs) Double bit flips, soft faults

Exascale computing (1000 PFlops/s)

- Larger core count: millions or even billions of cores
- Shorter Mean Time Between Failures (MTBF) μ



VANDERBILT
UNIVERSITY

Failures and Faults in HPC

Top ranked supercomputers in the US (June 2017)

Rank	Name	Laboratory	Technology	Cores	PFlops/s	MTBF
4	Titan	ORNL	Cray XK7	560,640	17.59	≈ 1 day
5	Sequoia	LLNL	BG/Q	1,572,864	17.17	≈ 1 day
6	Cori	LBNL	Cray XC40	622,336	14.01	≈ 1 day
9	Mira	ANL	BG/Q	786,432	8.59	≈ 1 day

Fail-stop errors Node failure, resource crashes

Silent errors or silent data corruptions (SDCs) Double bit flips, soft faults

Exascale computing (1000 PFlops/s)

- ▶ Larger core count: millions or even billions of cores
- ▶ Shorter Mean Time Between Failures (MTBF) μ

Coping with faults

- ▶ Build more reliable hardware!
- ▶ Make applications more fault tolerant!
- ▶ Design better resilience techniques/algorithms!



VANDERBILT
UNIVERSITY

Exascale platforms

► Hierarchical

- 10^5 or 10^6 nodes
- Each node equipped with 10^4 or 10^3 cores

► Failure-prone

MTBF – one node	1 year	10 years	120 years
MTBF – platform of 10^6 nodes	30sec	5mn	1h

More nodes → Shorter MTBF (Mean Time Between Failures)



VANDERBILT
UNIVERSITY

Exascale platforms

► Hierarchical

- 10^5 or 10^6 nodes
- Each node equipped with 10^4 or 10^3 cores

► Failure-prone

MTBF – one node	1 year	10 years	120 years
MTBF – platform of 10^6 nodes	30sec	5mn	1h

More nodes → Shorter MTBF (Mean Time Between Failures)

Exascale \neq Petascale x 1000 !!



VANDERBILT
UNIVERSITY

Even for today's platforms (courtesy F. Cappello)

Joint Laboratory for Petascale Computing

Also an issue at Petascale

INRIA NCSA


Fault tolerance becomes critical at Petascale (MTTI ≤ 1 day)
Poor fault tolerance design may lead to huge overhead

Overhead of checkpoint/restart

Cost of non optimal checkpoint intervals:

Today, 20% or more of the computing capacity in a large high-performance computing system is wasted due to failures and recoveries.


Dr. E.N. (Mootaz) Elnozahy et al. *System Resilience at Extreme Scale, DARPA*



Classic Resilience Techniques for HPC

Fail-stop errors (instantaneous error detection)

Standard approach periodic checkpointing, rollback and recovery





VANDERBILT
UNIVERSITY

Classic Resilience Techniques for HPC

Fail-stop errors (instantaneous error detection)

Standard approach periodic checkpointing, rollback and recovery



Without optimizations, currently C/R could take up to one hour !



VANDERBILT
UNIVERSITY

Sources of failures

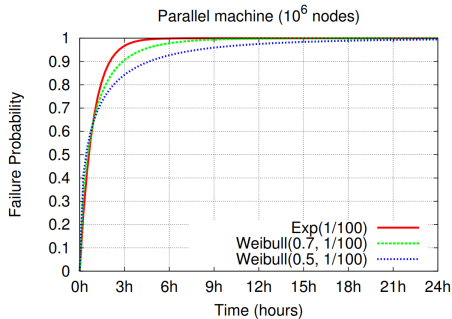
- ▶ Many types of faults: software error, hardware malfunction, memory corruption
 - ▶ Software errors: Applications, OS bug (kernel panic), communication libs, File system error, etc
 - ▶ Hardware errors: Disks, processors, memory, network
- ▶ Many possible behaviors: silent, transient, unrecoverable
- ▶ **Restrict to faults that lead to application failures**
 - ▶ This includes most hardware faults, and some software ones
- ▶ Will use terms fault and failure interchangeably
- ▶ **Silent errors (SDC) addressed later in the lecture**



VANDERBILT
UNIVERSITY

Values from real systems

- ▶ MTBF of one processor: between 1 and 125 years
- ▶ MTBF for systems: between several hours to one day



Failure trace archive from INRIA (<http://fta.inria.fr>)

Computer Failure Data Repository from LANL (<http://institutes.lanl.gov/data/fdata>)



Process Checkpointing

Goal

Save the current state of the process

- ▶ FT Protocols save a possible state of the parallel application

Two techniques

- ▶ User-level checkpointing
- ▶ System-level checkpointing

- ▶ Blocking call
- ▶ Asynchronous call



VANDERBILT
UNIVERSITY

User-level checkpointing

User code serializes the state of the process in a file

- ▶ Usually small(er than system-level checkpointing)
- ▶ Portability
- ▶ Diversity of use

- ▶ Hard to implement if preemptive checkpointing is needed
- ▶ Loss of the functions call stack
 - ▶ code full of jumps
 - ▶ loss of internal library state

FTI (Fault Tolerance Interface)

Multi-level checkpointing tool for MPI applications



VANDERBILT
UNIVERSITY

System-level checkpointing

- ▶ Different possible implementations: OS syscall; dynamic library; compiler assisted
- ▶ Create a serial file that can be loaded in a process image. Usually on the same architecture, same OS, same software environment.
- ▶ **Entirely transparent**
- ▶ Preemptive (often needed for library-level checkpointing)
- ▶ Lack of portability
- ▶ Large size of checkpoint (\approx memory footprint)

BLCR

Requires linking with the code

CRIU

User-space tool that can freeze a running container (or an individual application) and checkpoint its state to disk



VANDERBILT
UNIVERSITY

Blocking / Asynchronous call

Blocking Checkpointing

- ▶ Relatively intuitive: `checkpoint(filename)`
- ▶ Cost: no process activity during the whole checkpoint operation.
- ▶ Can be linear in the size of memory and in the size of modified files

Asynchronous Checkpointing

- ▶ System-level approach: make use of copy on write or fork syscall
- ▶ User-level approach: critical sections, when needed



VANDERBILT
UNIVERSITY

Remote Reliable Storage

- ▶ Intuitive. I/O intensive. Disk usage.

Memory Hierarchy

- ▶ local memory
- ▶ local disk (SSD, HDD)
- ▶ remote disk
 - ▶ Scalable Checkpoint Restart Library
 - ▶ FTI

Checkpoint is valid when finished on reliable storage



VANDERBILT
UNIVERSITY

How to use checkpointing for your application?

1 Choose the checkpointing tool that fits your application

- ▶ Do you have access to the code of the application?
- ▶ What is the memory footprint?
- ▶ How much time do you want to invest in changing your application?
- ▶ What overhead is acceptable?
- ▶ Is there a plugin in SLURM capable of taking checkpoints for you?

2 Choose the storage space for your application

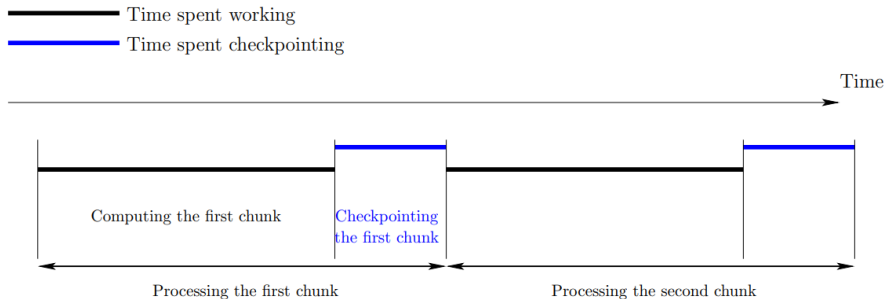
- ▶ Permanent storage or local memory

3 Choose how often you take checkpoints



VANDERBILT
UNIVERSITY


Checkpointing cost



Blocking model: while a checkpoint is taken, no computation can be performed



Checkpointing cost



Time lost

- ▶ Downtime required to fix the error
- ▶ Time it takes to restart the application
- ▶ Average lost time is half the time between two checkpoints
- ▶ In addition to the time to take the checkpoints



Checkpointing interval

Optimization problem

- ▶ Studied by multiple research projects
 - ▶ Failure rate of different hardware components
 - ▶ Application behavior
- ▶ **Daly's optimal checkpoint sequence from 2006 is still valid**
 $\sqrt{2\mu C}$



VANDERBILT
UNIVERSITY

Lesson learned for fail-stop failures

- ▶ Tsubame 2: 962 failures during last 18 months so $\mu = 13$ hrs
- ▶ Blue Waters: 2-3 node failures per day
- ▶ Titan: a few failures per day
- ▶ Tianhe 2: wouldn't say

$$T_{\text{opt}} = \sqrt{2\mu C} \Rightarrow \text{WASTE}[\text{opt}] \approx \sqrt{\frac{2C}{\mu}}$$

Petascale:	$C = 20$ min	$\mu = 24$ hrs	$\Rightarrow \text{WASTE}[\text{opt}] = 17\%$
Scale by 10:	$C = 20$ min	$\mu = 2.4$ hrs	$\Rightarrow \text{WASTE}[\text{opt}] = 53\%$
Scale by 100:	$C = 20$ min	$\mu = 0.24$ hrs	$\Rightarrow \text{WASTE}[\text{opt}] = 100\%$



Table of content

- ▶ Failures and faults in HPC
- ▶ Checkpointing
 - ▶ Theory
 - ▶ System and application level solutions currently used
- ▶ **Fault tolerant MPI**
- ▶ Designing a resilient application
- ▶ Silent errors



VANDERBILT
UNIVERSITY

Message Passing Interface Most popular middleware for multi-node programming in HPC

MPI Standard 3.0, p. 21, l. 24:25

This document does not specify the state of a computation after an erroneous MPI call has occurred.

Open MPI (<http://www.open-mpi.org>) **MPICH** (<http://www.mcs.anl.gov/mpi/mpich/>)

- ▶ Default: on failure detection, the runtime kills all processes. Can be de-activated by a runtime switch
- ▶ Errors might be reported to MPI processes. In that case MPI might be partly usable.



VANDERBILT
UNIVERSITY

MPI-Next-FT proposal: ULFM

Goal

Resume Communication Capability for MPI (and nothing more)

- ▶ Error reporting indicates impossibility to carry an operation
 - ▶ State of MPI is unchanged for operations that can continue (i.e. if they do not involve a dead process)
- ▶ Errors are non uniformly returned
 - ▶ (Otherwise, synchronizing semantic is altered drastically with high performance impact)

New APIs

- ▶ REVOKE allows to resolve non-uniform error status
- ▶ SHRINK allows to rebuild error-free communicators




VANDERBILT
UNIVERSITY

MPI-Next-FT proposal: ULFM

Errors are visible only for operations that cannot complete

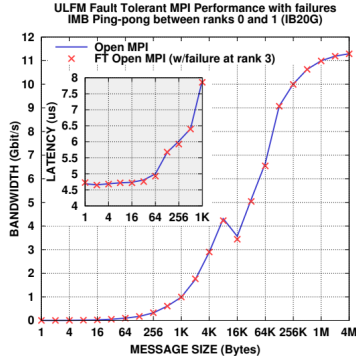

Error Reporting

- ▶ Operations that cannot complete return `ERR_PROC_FAILED`, or `ERR_PENDING` if appropriate
- ▶ State of MPI Objects is unchanged (communicators etc.)
- ▶ Operations that can be completed return `MPI_SUCCESS` point to point operations between non-failed ranks can continue



VANDERBILT
UNIVERSITY

MPI-Next-FT proposal: ULFM



OpenMPI support

Branch of Open MPI (www.open-mpi.org)



VANDERBILT
UNIVERSITY

Table of content

- ▶ Failures and faults in HPC
- ▶ Checkpointing
 - ▶ Theory
 - ▶ System and application level solutions currently used
- ▶ Fault tolerant MPI
- ▶ **Designing a resilient application**
- ▶ Silent errors



VANDERBILT
UNIVERSITY

Application level resiliency

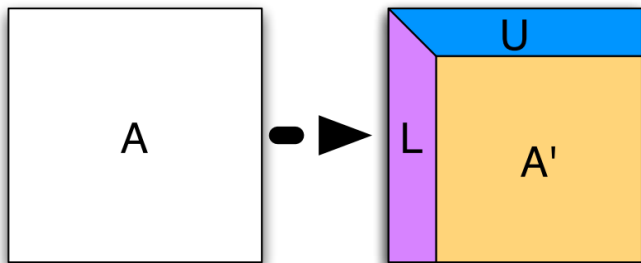
Applications can build customized methods for tolerating failures

- ▶ Depending on the application behavior some can mask failures all together
 - ▶ Neural networks
 - ▶ Heat equation
- ▶ **Algorithm based fault tolerance**
 - ▶ ABFT for Linear Algebra applications



VANDERBILT
UNIVERSITY

Example: block LU/QR factorization




- ▶ Solve $A \cdot x = b$ (hard)
- ▶ Transform A into a LU factorization
- ▶ Solve $L \cdot y = B \cdot b$, then $U \cdot x = y$



Example: block LU/QR factorization

TRSM - Update row block



GETF2: factorize a column block
GEMM: Update the trailing matrix

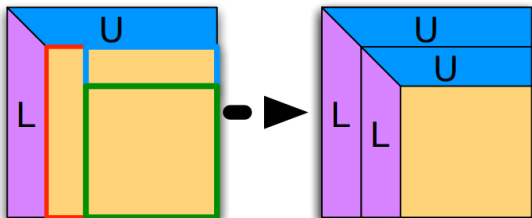
- ▶ Solve $A \cdot x = b$ (hard)
- ▶ Transform A into a LU factorization
- ▶ Solve $L \cdot y = B \cdot b$, then $U \cdot x = y$



VANDERBILT
UNIVERSITY

Example: block LU/QR factorization

TRSM - Update row block



GETF2: factorize a column block
GEMM: Update the trailing matrix

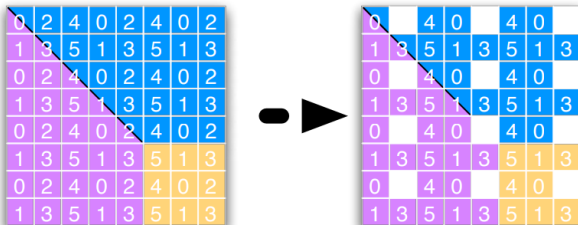
- ▶ Solve $A \cdot x = b$ (hard)
- ▶ Transform A into a LU factorization
- ▶ Solve $L \cdot y = B \cdot b$, then $U \cdot x = y$



VANDERBILT
UNIVERSITY




Example: block LU/QR factorization

Failure of rank 2



- ▶ 2D Block Cyclic Distribution (here 2 x 3)
- ▶ A single failure → many data elements will be lost

How do checksum work?




- For vector matrix multiplication, compute row/column wise checksum
- For LU factorization naive implementations will fail



VANDERBILT
UNIVERSITY


Checksum for LU



- Replication is needed




Algorithm Based Fault Tolerant QR decomposition



- Checksum: invertible operation on the data of the row / column
- Checksum blocks are doubled, to allow recovery when data and checksum are lost together

Algorithm Based Fault Tolerant QR decomposition




- ▶ Checksum: invertible operation on the data of the row / column
- ▶ Checksum replication can be avoided by dedicating computing resources



VANDERBILT
UNIVERSITY

Algorithm Based Fault Tolerant QR decomposition




Idea of ABFT

Applying the operation on data and checksum preserves the checksum properties



VANDERBILT
UNIVERSITY

Algorithm Based Fault Tolerant QR decomposition




Idea of ABFT

For the part of the data that is not updated this way, the checksum must be re-calculated



VANDERBILT
UNIVERSITY

In case of a failure




- In case of failure, conclude the operation, then
 - Missing Data = Checksum - Sum(Existing Data) s



VANDERBILT
UNIVERSITY

In case of a failure




- ▶ In case of failure, conclude the operation, then
 - ▶ Missing Data = Checksum - Sum(Existing Data) s



VANDERBILT
UNIVERSITY

ABFT QR decomposition: performance



Open MPI; Kraken supercomputer;



VANDERBILT
UNIVERSITY

Table of content

- ▶ Failures and faults in HPC
- ▶ Checkpointing
 - ▶ Theory
 - ▶ System and application level solutions currently used
- ▶ Fault tolerant MPI
- ▶ Designing a resilient application
- ▶ **Silent errors**



VANDERBILT
UNIVERSITY

What are silent failures?

- ▶ Instantaneous error detection → fail-stop failures, e.g. resource crash
- ▶ Silent errors (data corruption) → detection latency

Silent error detected only when the corrupt data is activated

- ▶ Includes some software faults, some hardware errors (soft errors in L1 cache), double bit flip
- ▶ Cannot always be corrected by ECC memory



VANDERBILT
UNIVERSITY

Should we be afraid? (courtesy AI Geist)

Fear of the Unknown

Hard errors – permanent component failure either HW or SW
(hung or crash)

Transient errors – a blip or short term failure of either HW or SW

Silent errors – undetected errors either hard or soft, due to lack of detectors for a component or inability to detect (transient effect too short). Real danger is that answer may be incorrect but the user wouldn't know.

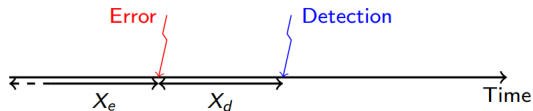
**Statistically, silent error rates are increasing.
Are they really? Its fear of the unknown**

Are silent errors really a problem
or just monsters under our bed?



VANDERBILT
UNIVERSITY

Coupling checkpointing with verification




Error and detection latency

- ▶ Last checkpoint may have saved an already corrupted state
- ▶ Possible solution: Saving k checkpoints (Lu, Zheng and Chien):
 - ❶ Critical failure when all live checkpoints are invalid



VANDERBILT
UNIVERSITY

Coupling checkpointing with verification




Error and detection latency

- ▶ Last checkpoint may have saved an already corrupted state
- ▶ Possible solution: Saving k checkpoints (Lu, Zheng and Chien):
 - ❶ Critical failure when all live checkpoints are invalid
Assume unlimited storage resources
 - ❷ Which checkpoint to roll back to?



VANDERBILT
UNIVERSITY

Coupling checkpointing with verification




Error and detection latency

- ▶ Last checkpoint may have saved an already corrupted state
- ▶ Possible solution: Saving k checkpoints (Lu, Zheng and Chien):
 - ❶ Critical failure when all live checkpoints are invalid
Assume unlimited storage resources
 - ❷ Which checkpoint to roll back to?
Assume verification mechanism



VANDERBILT
UNIVERSITY

Coupling checkpointing with verification



- ▶ The most promising approach: checkpointing + verification (error detection)
- ▶ Use verification to correct some errors (ABFT)

A lot of unknowns

- ▶ Error rate? MTBE?
- ▶ Selective reliability?
- ▶ New algorithms beyond iterative? matrix-product, FFT



- ▶ **Multiple approaches to Fault Tolerance**
- ▶ **Application-Specific Fault Tolerance will always provide more benefits:**
 - ▶ Checkpoint Size Reduction (when needed)
 - ▶ Portability (can run on different hardware, different deployment, etc..)
 - ▶ Diversity of use (can be used to restart the execution and change parameters in the middle)
- ▶ **General Purpose Fault Tolerance is a required feature of the platforms**
 - ▶ Not every computer scientist needs to learn how to write fault-tolerant applications
- ▶ **Requirements of a more Fault-friendly programming environment**
 - ▶ MPI-Next evolution
 - ▶ Other programming environments?



Fault-Tolerance Techniques for High-Performance Computing

Springer Book, Computer Communications and Networks series, 2015

Editors: Thomas Herault and Yves Robert

Checkpointing tools

- ▶ FTI: Fault Tolerance Interface, application level checkpointing using multi-level storage, <https://github.com/leobago/fti>
- ▶ BLCR: Berkeley Lab Checkpoint/Restart for Linux (system level checkpointing tool), <https://crd.lbl.gov/departments/computer-science/CLaSS/research/BLCR/>
- ▶ CRIU: Checkpoint and Restore in Userspace, system level tool to checkpoint/restore Linux tasks, <https://github.com/checkpoint-restore/criu>

More info on the ABFT example

Algorithm-based fault tolerance applied to high performance computing,
Bosilca G. et al., JPDC 69, 4 (2009)



VANDERBILT
UNIVERSITY