

# Solving Linear Equations, FFT and Sorting on Distributed-Memory Architectures

SC3260/5260 High-Performance Computing

Hongyang Sun

([hongyang.sun@vanderbilt.edu](mailto:hongyang.sun@vanderbilt.edu))

Vanderbilt University

Spring 2020



# Solving System of Linear Equations

- Solving  $Ax = b$  as a benchmark to rank world's top supercomputers (<https://www.top500.org/>).
- We will discuss the classical two-step approach:
  - 1) Gaussian Elimination;
  - 2) Back Substitution.
- We assume a 2D block data distribution on  $\sqrt{P} \times \sqrt{P}$  procs.
  - Proc  $P_{ij}$  holds matrix block  $A_{ij}$
  - Diagonal processor  $P_{ii}$  holds vector block  $b_i$  and computes  $x_i$

$P_{00}$ $A_{00}, b_0$ $\rightarrow x_0$	$P_{01}$ $A_{01}$	$P_{02}$ $A_{02}$
$P_{10}$ $A_{10}$	$P_{11}$ $A_{11}, b_1$ $\rightarrow x_1$	$P_{12}$ $A_{12}$
$P_{20}$ $A_{20}$	$P_{21}$ $A_{21}$	$P_{22}$ $A_{22}, b_2$ $\rightarrow x_2$

# A Simple Example

$$\blacksquare \begin{cases} 2x_0 + 4x_1 + 6x_2 = 8 \\ 3x_0 + 9x_1 + 15x_2 = 21 \\ 5x_0 + 12x_1 + 7x_2 = 2 \end{cases} \quad \rightarrow \quad \begin{bmatrix} 2 & 4 & 6 \\ 3 & 9 & 15 \\ 5 & 12 & 7 \end{bmatrix} \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 8 \\ 21 \\ 2 \end{bmatrix}$$

*System of linear equations*

*Matrix form  $Ax = b$*

# A Simple Example

## ■ **Step 1: Gaussian Elimination**

- Reduce  $Ax = b$  to  $Ux = y$  (*row echelon form*), where  $U$  is an upper triangular matrix (with 1 on diagonal).
- Achieved by dividing a row by a non-zero number and subtracting a row by a multiple of another row.

$$\begin{array}{ccc} \left[ \begin{array}{ccc|c} 2 & 4 & 6 & 8 \\ 3 & 9 & 15 & 21 \\ 5 & 12 & 7 & 2 \end{array} \right] & \rightarrow & \left[ \begin{array}{ccc|c} 1 & 2 & 3 & 4 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 2 \end{array} \right] \\ A & & U \quad y \end{array}$$

A lower triangular matrix  $L$  records intermediate steps of this process

$$\begin{array}{ccc} & \searrow & \\ \left[ \begin{array}{ccc} 2 & 0 & 0 \\ 3 & 3 & 0 \\ 5 & 2 & -12 \end{array} \right] & & \\ & L & \end{array}$$

# A Simple Example

## ■ **Step 2: Back Substitution**

- Solve for  $x$  in  $Ux = y$  in reverse order from  $x_{n-1}$  to  $x_0$

$$\begin{array}{ccc} \left[ \begin{array}{ccc|c} 1 & 2 & 3 & 4 \\ 0 & 1 & 2 & 3 \\ 0 & 0 & 1 & 2 \end{array} \right] & \rightarrow & \begin{cases} x_0 + 2x_1 + 3x_2 = 4 \\ x_1 + 2x_2 = 3 \\ x_2 = 2 \end{cases} \rightarrow \begin{cases} x_0 = 0 \\ x_1 = -1 \\ x_2 = 2 \end{cases} \\ U & & y \end{array}$$

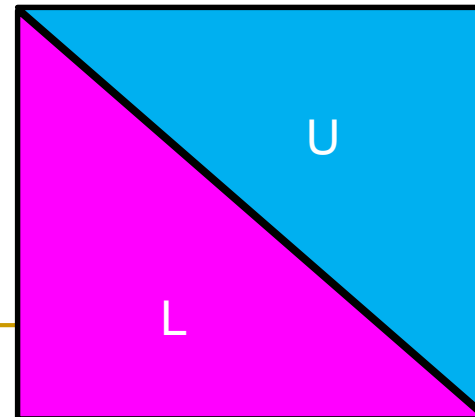
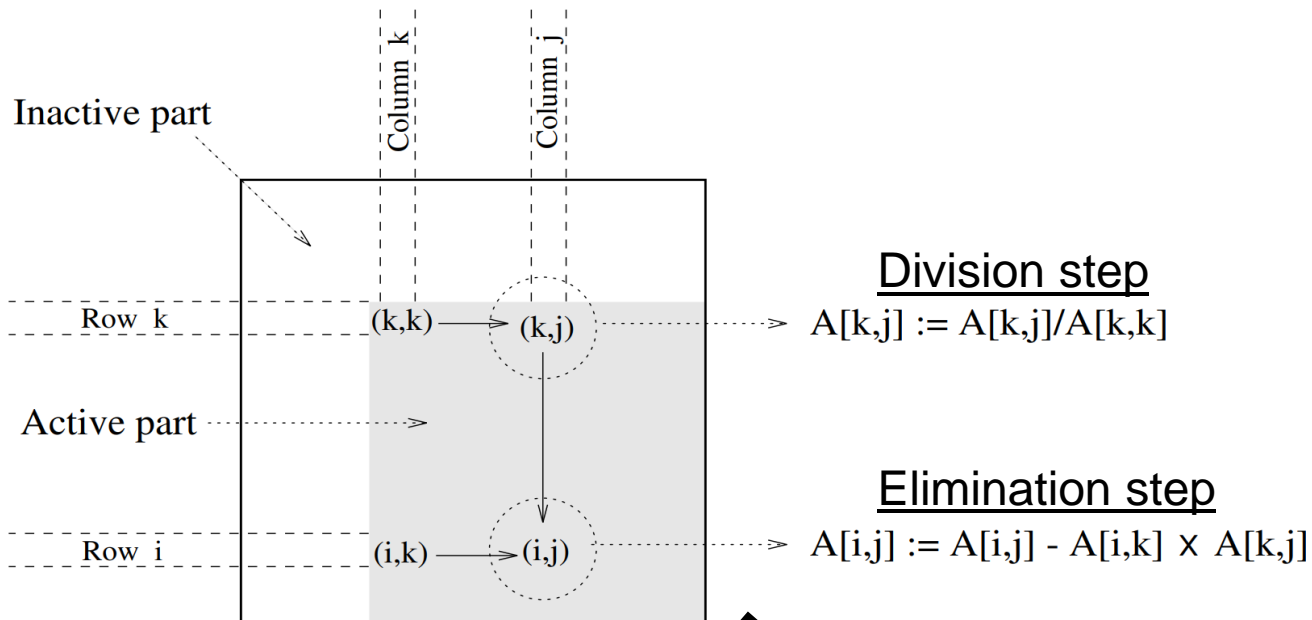
*For an  $n \times n$  matrix  $A$ , Step 1 (Gaussian Elimination) is more expensive taking  $O(n^3)$  time, and Step 2 (Back Substitution) takes only  $O(n^2)$  time.*

# Pseudocodes of Two Steps

```
1.  procedure GAUSSIAN_ELIMINATION ( $A, b, y$ )
2.  begin
3.    for  $k := 0$  to  $n - 1$  do           /* Outer loop */
4.      begin
5.        for  $j := k + 1$  to  $n - 1$  do
6.           $A[k, j] := A[k, j] / A[k, k];$  /* Division step */
7.           $y[k] := b[k] / A[k, k];$ 
8.           $A[k, k] := 1;$ 
9.          for  $i := k + 1$  to  $n - 1$  do
10.           begin
11.             for  $j := k + 1$  to  $n - 1$  do
12.                $A[i, j] := A[i, j] - A[i, k] \times A[k, j];$  /* Elimination step */
13.                $b[i] := b[i] - A[i, k] \times y[k];$ 
14.                $A[i, k] := 0;$ 
15.             endfor;           /* Line 9 */
16.           endfor;           /* Line 3 */
17.        end GAUSSIAN_ELIMINATION
```

```
1.  procedure BACK_SUBSTITUTION ( $U, x, y$ )
2.  begin
3.    for  $k := n - 1$  downto  $0$  do /* Main loop */
4.      begin
5.         $x[k] := y[k];$ 
6.        for  $i := k - 1$  downto  $0$  do
7.           $y[i] := y[i] - x[k] \times U[i, k];$ 
8.        endfor;
9.      end BACK_SUBSTITUTION
```

# Illustration of Gaussian Elimination



# Gaussian Elimination in 2D Blocks

- For step  $k = 0, 1, \dots, \sqrt{P} - 1$ :
  - a) Processor  $P_{kk}$  performs local division and elimination on matrix block  $A_{kk}$  (resulting in  $L_{kk}$  and  $U_{kk}$ );
  - b) Processor  $P_{kk}$  broadcasts  $L_{kk}$  to all processors  $P_{kj}$  ( $j > k$ ), and  $U_{kk}$  to all processors  $P_{ik}$  ( $i > k$ );
  - c) Processors  $P_{kj}$  and  $P_{ik}$  ( $i, j > k$ ) perform local division and elimination (resulting in  $U_{kj}$  and  $L_{ik}$ , respectively);
  - d) Processors  $P_{kj}$  broadcasts  $U_{kj}$  down to other processors in same column, and processors  $P_{ik}$  broadcasts  $L_{ik}$  to other processors to the right in same row;
  - e) Processors  $P_{ij}$  ( $i, j > k$ ) perform local division and elimination (resulting in modified block  $\widetilde{A}_{ij}$ ).

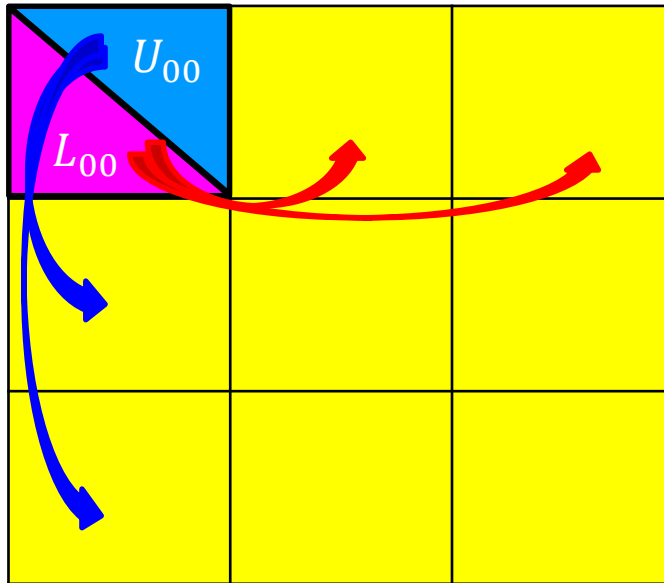
---

*Next slide illustrates one step for  $P = 9$  processors* 8

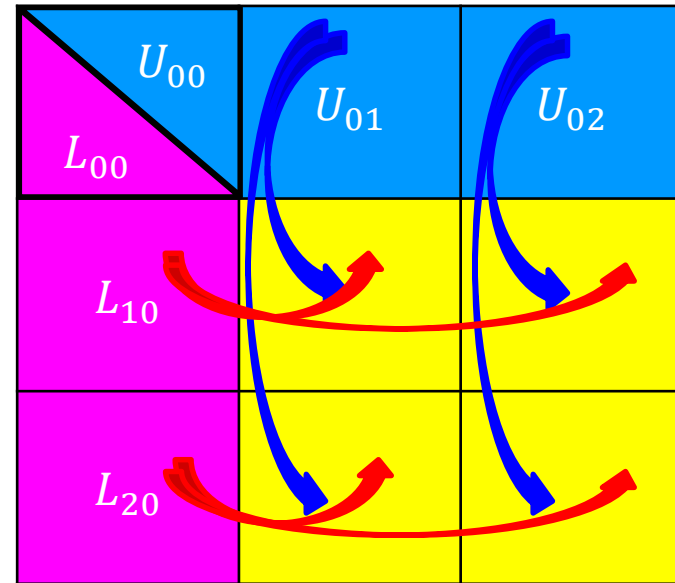


# Gaussian Elimination in 2D Blocks

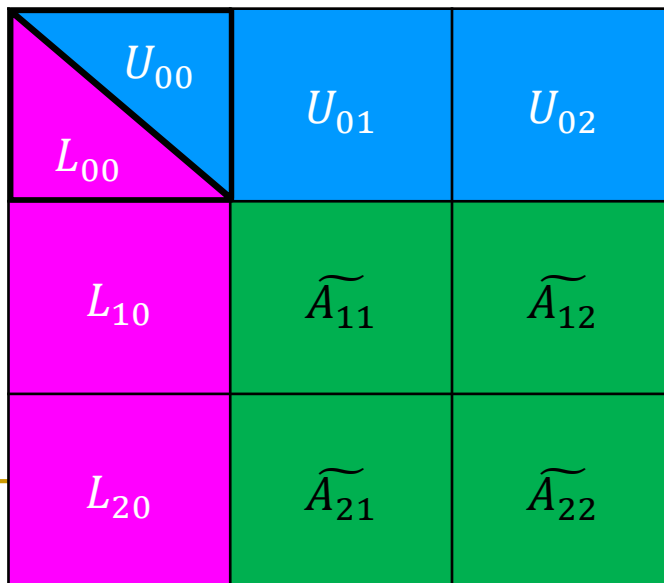
Step  $k = 0$   
(a)(b)



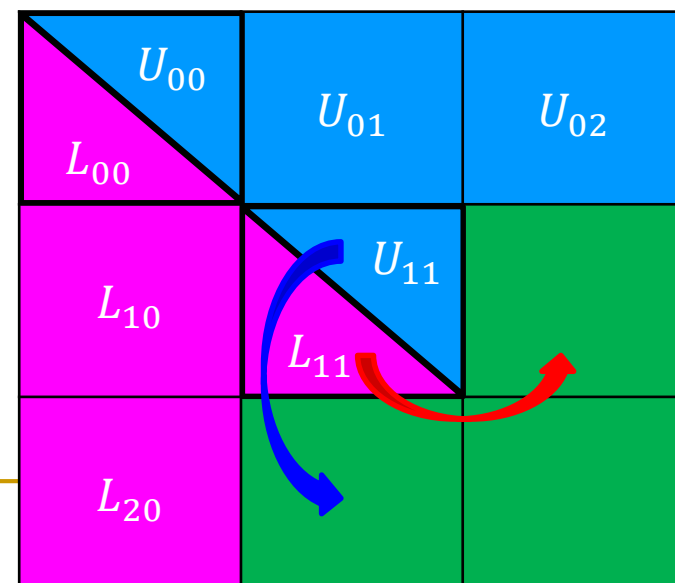
Step  $k = 0$   
(c)(d)



Step  $k = 0$   
(e)



Step  $k = 1$   
(a)(b)

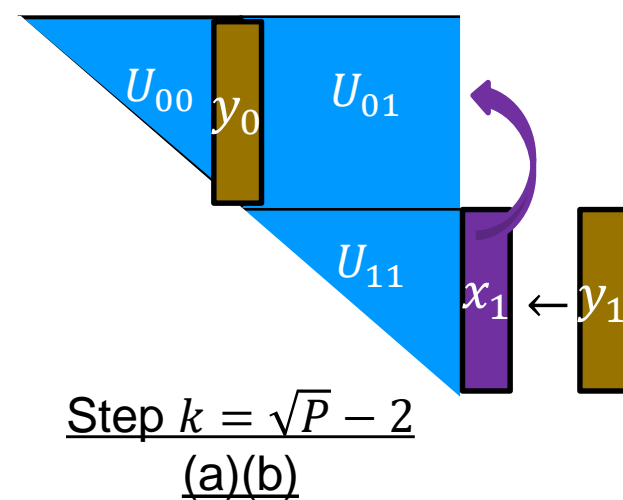
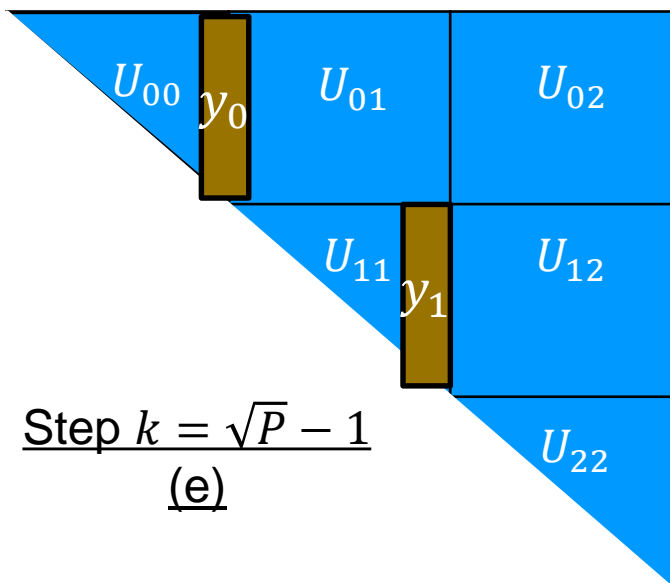
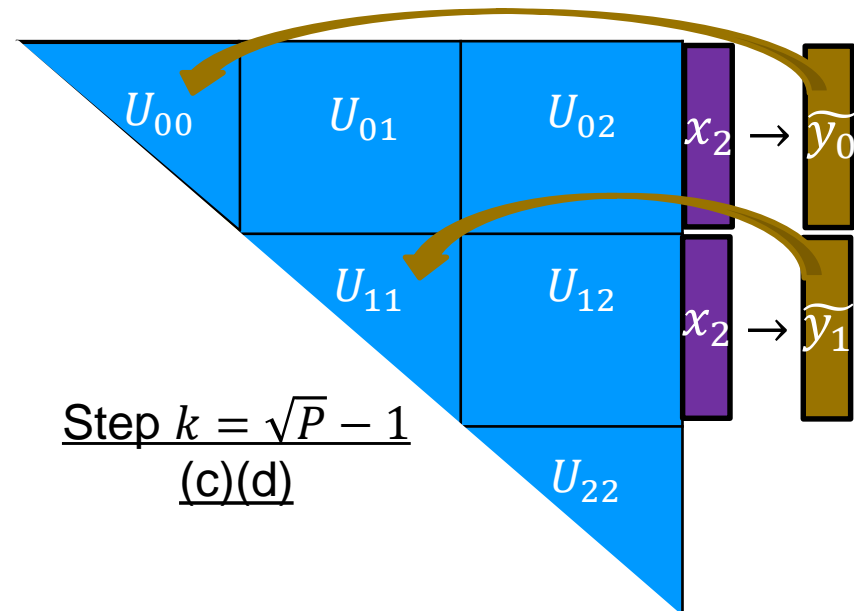
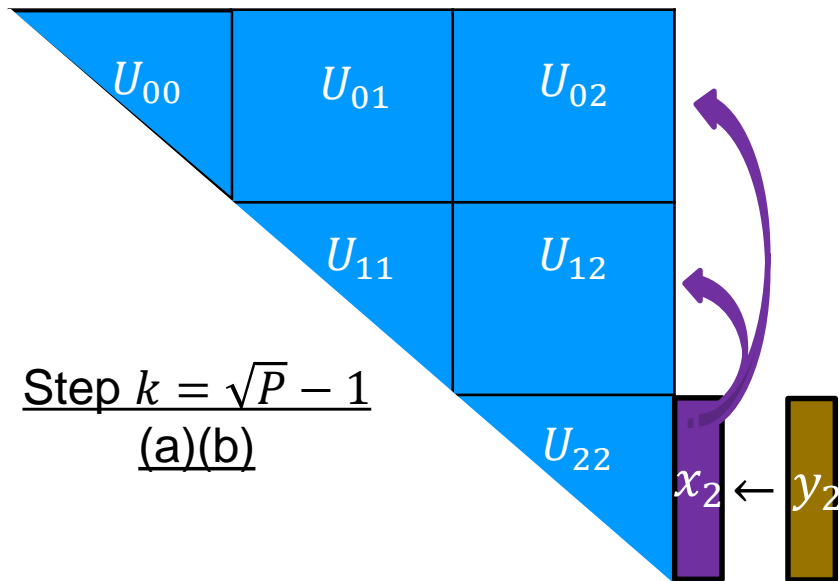


# Back Substitution in 2D Blocks

- For step  $k = \sqrt{P} - 1, \sqrt{P} - 2, \dots, 0$ :
  - a) Processor  $P_{kk}$  performs local back substitution on matrix block  $U_{kk}$  and  $y_k$  (resulting in  $x_k$ );
  - b) Processor  $P_{kk}$  broadcasts  $x_k$  to all processors  $P_{jk}$  ( $j < k$ );
  - c) Processors  $P_{jk}$  ( $j < k$ ) computes  $\tilde{y}_j \leftarrow U_{jk} \times x_k$ ;
  - d) Processors  $P_{jk}$  ( $j < k$ ) sends  $\tilde{y}_j$  to processors  $P_{j,j}$ ;
  - e) Processors  $P_{j,j}$  ( $j < k$ ) updates  $y_j \leftarrow y_j - \tilde{y}_j$ .

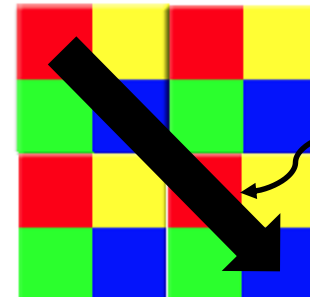
*Next slide illustrates one step for  $P = 9$  processors*

# Back Substitution in 2D Blocks

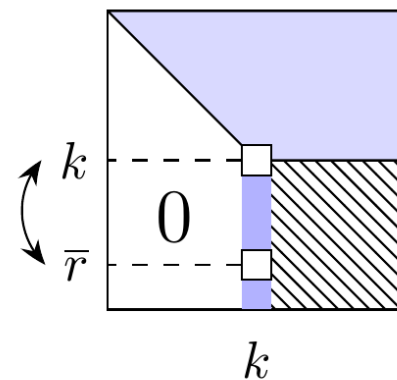


# Solving $Ax = b$ in Practice

- **2D Block-cyclic** distribution is used (for better load balance);
- **Partial pivoting** is used during Gaussian Elimination (for numerical stability);
- **LU factorization** of matrix  $A$  (for reusing left side).



Order of computation for Gaussian Elimination



Partial pivoting by row permutation

*These and other optimizations are implemented in the software package High-Performance Linpack (HPL) (<http://www.netlib.org/benchmark/hpl/>).*

# Fast Fourier Transform (FFT)

- **Discrete Fourier Transform (DFT):** converts a discrete signal of  $N$  samples from time domain  $(x_0, x_1, \dots, x_{N-1})$  to frequency domain  $(y_0, y_1, \dots, y_{N-1})$ , or vice versa:

$$y_j = \sum_{k=0}^{N-1} w_N^{jk} \cdot x_k, \forall j = 0, 1, \dots, N-1$$

where  $w_N = e^{-\frac{2\pi i}{N}}$  is the  $N$ -th root of unity, called the *twiddle factor*.

- **Fast Fourier Transform (FFT):** any algorithm that computes DFT in  $O(n \log n)$  time instead of the naïve  $O(n^2)$  complexity.

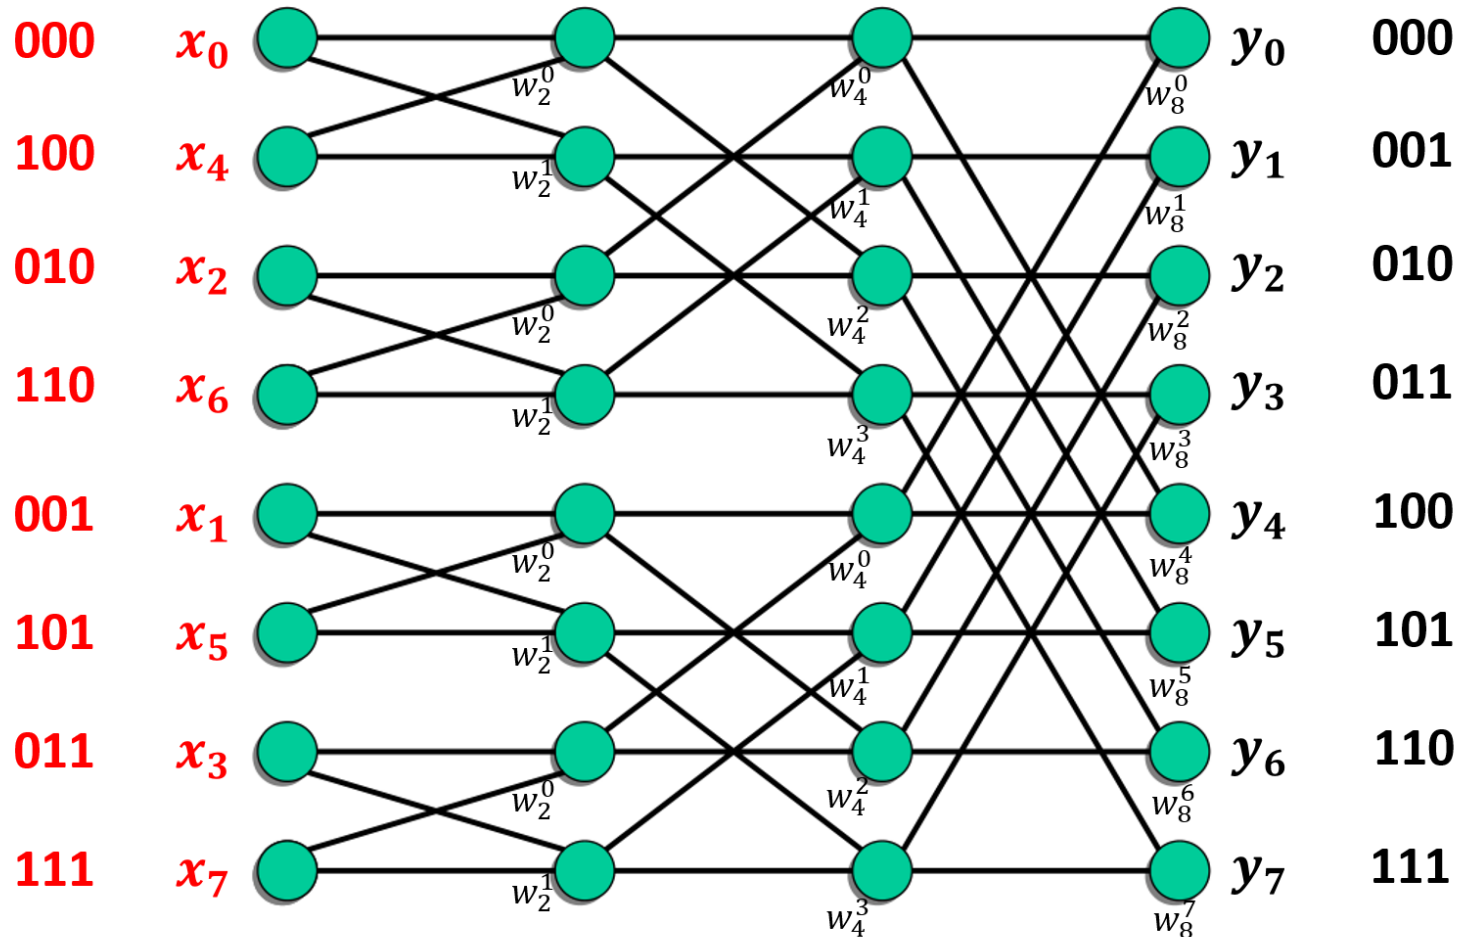
# Fast Fourier Transform (FFT)

- FFT has numerous applications in science and engineering, and is considered to be one of the most important algorithms in the 20<sup>th</sup> century!
- We consider the classical *radix-2* algorithm and study two parallel implementations:
  - 1) **Binary-Exchange Algorithm;**
  - 2) **Transpose Algorithm.**
- For simplicity, we assume that both  $N$  and  $P$  (number of processors) are powers-of-2, i.e.,  $N = 2^r$ ,  $P = 2^d$  and  $d \leq r$ .

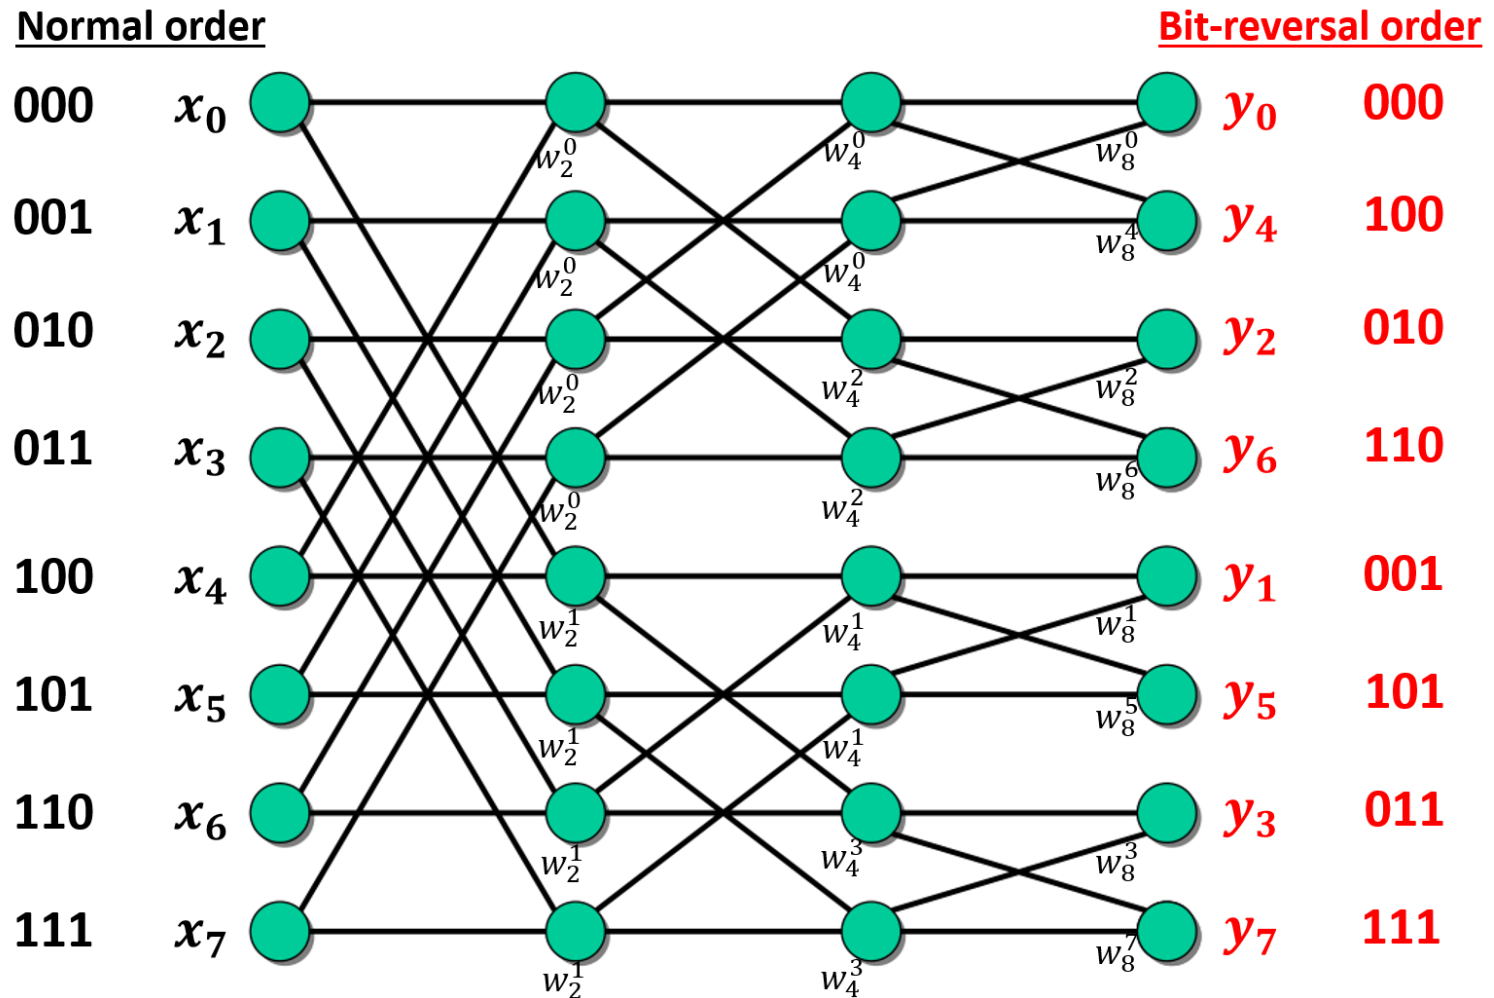
# Radix-2 FFT (Decimation in Time)

Bit-reversal order

Normal order



# Radix-2 FFT (Decimation in Frequency)

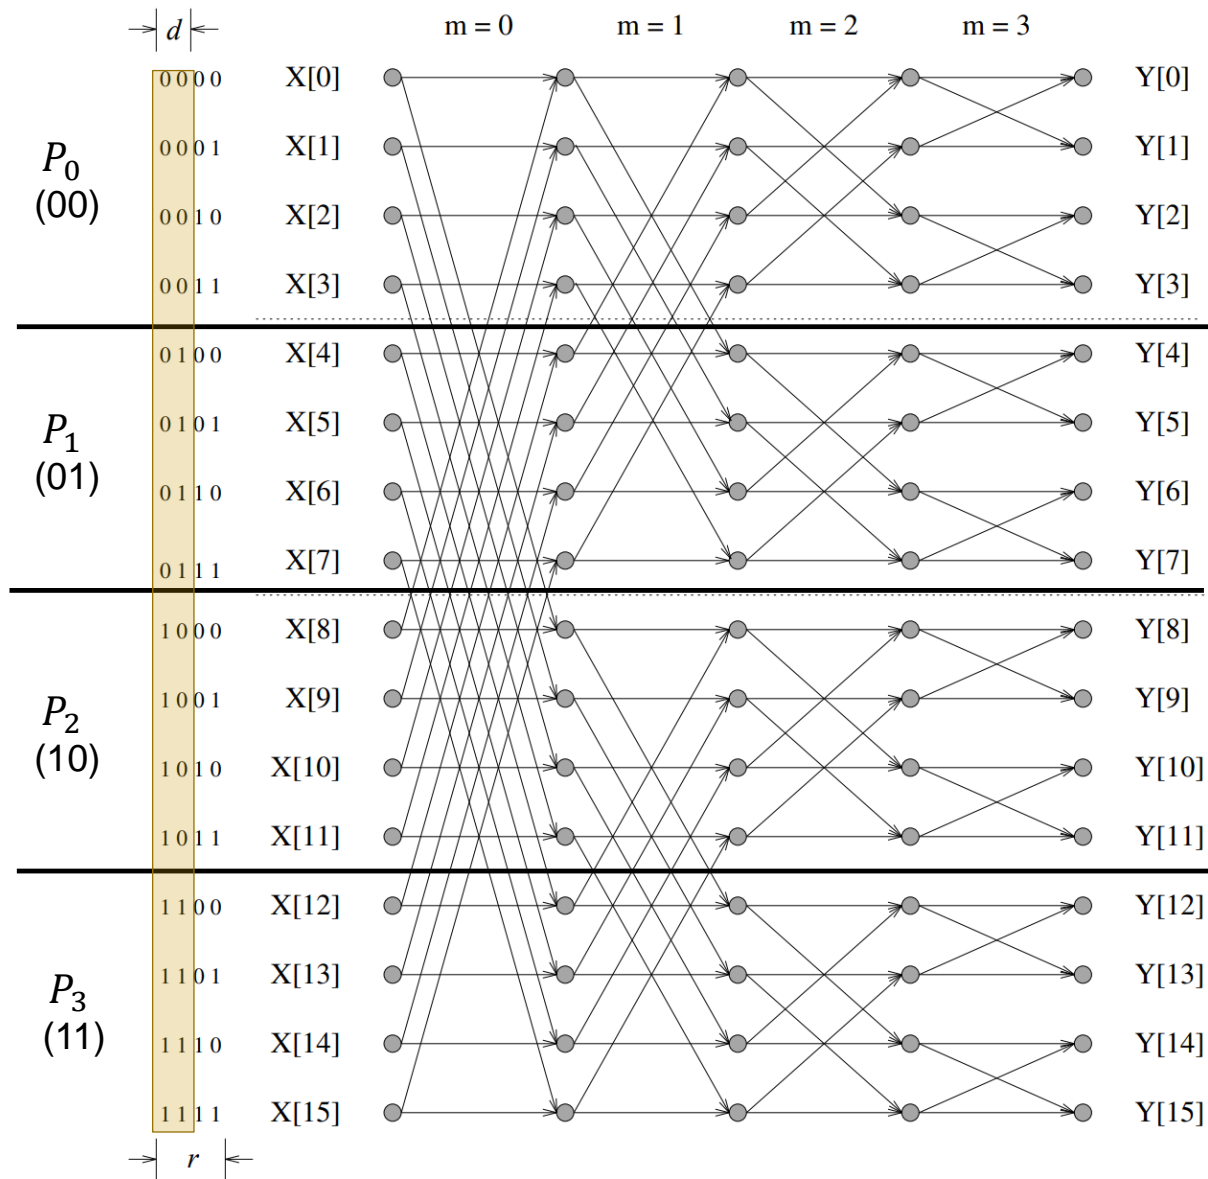




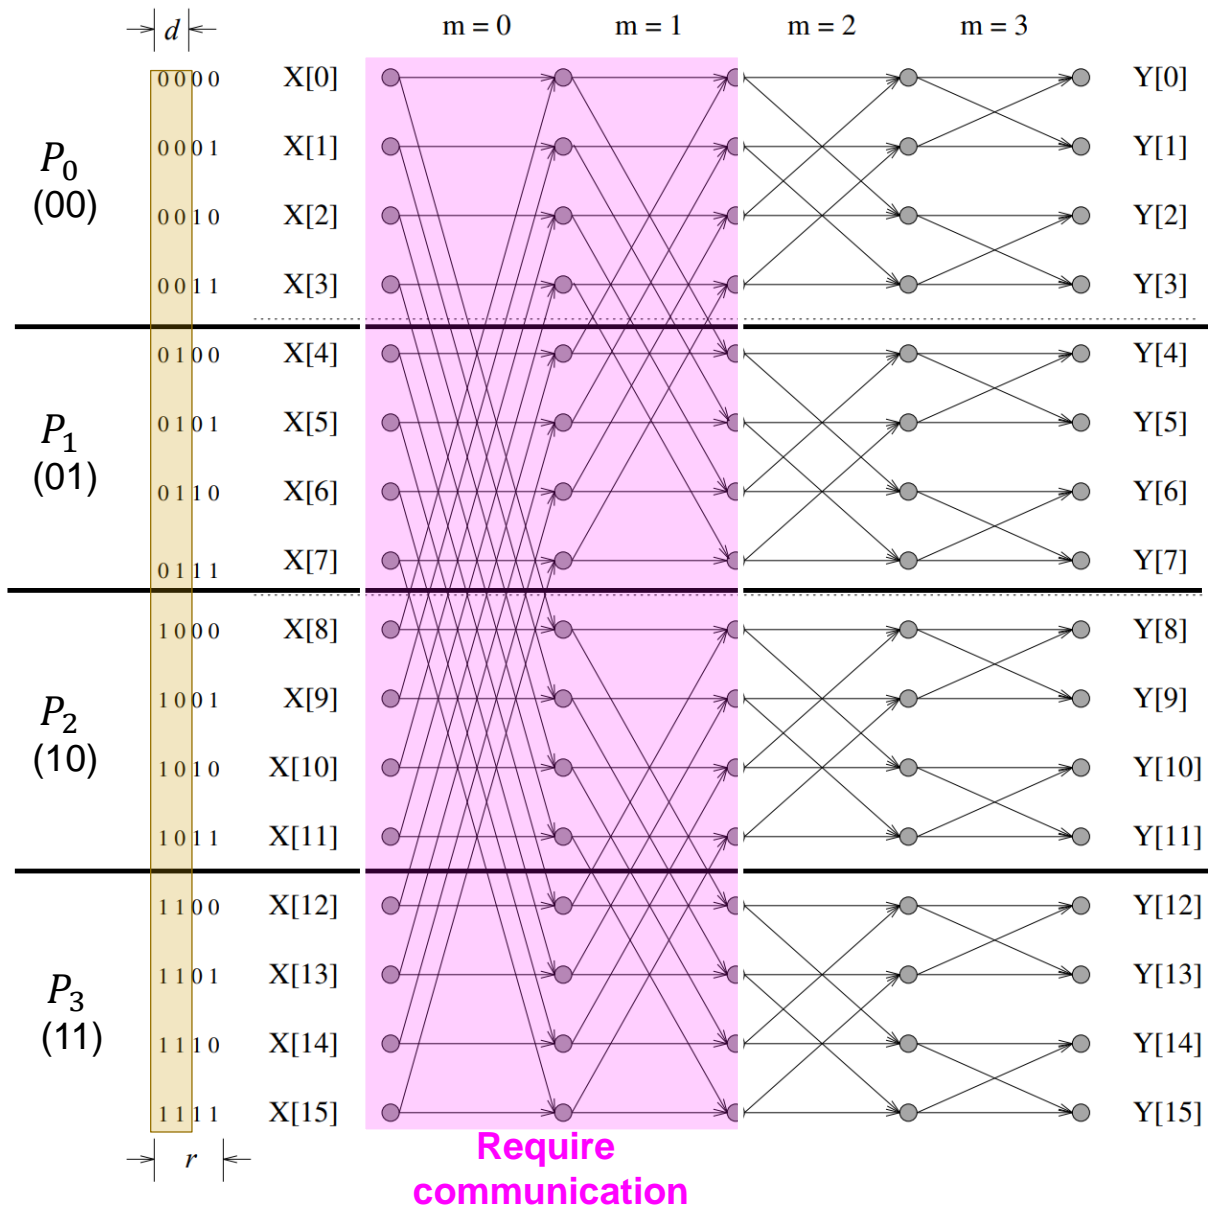
# (1) Binary-Exchange Algorithm

- We assume **1D block** data distribution for the DIF implementation of the radix-2 algorithm:
  - Each processor holds  $N/P$  consecutive input values in  $x$  and responsible for computing the corresponding output values in  $y$ .
- The algorithm works in two phases:
  - Phase 1 (w/ communication): step  $m = 0, 1, \dots, d - 1$ :
    - a) Each processor  $P_i$  exchanges  $N/P$  numbers with processor  $P_{i'}$ , where  $P_i$  and  $P_{i'}$  only differ in the  $m$ -th bit (from left);
    - b) Each processor  $P_i$  does  $N/P$  local complex computations;
  - Phase 2 (local computation): step  $m = d, \dots, r - 1$ :
    - a) Each processor  $P_i$  does  $N/P$  local complex computations;

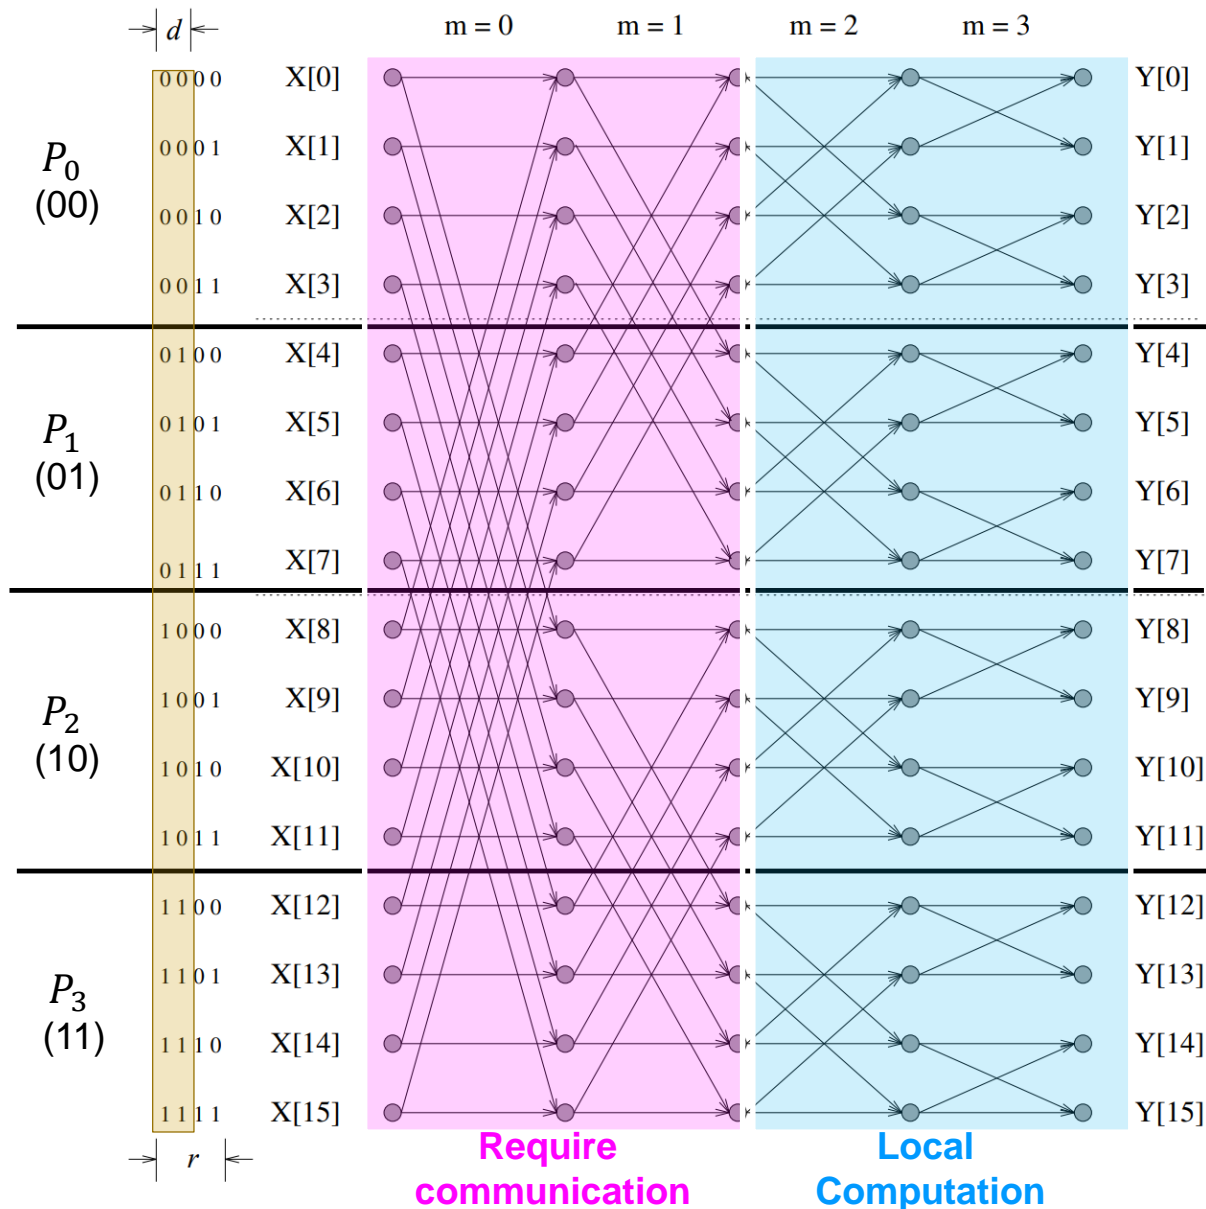
# (1) Binary-Exchange Algorithm



## (1) Binary-Exchange Algorithm



# (1) Binary-Exchange Algorithm



## (2) Transpose Algorithm


- We assume:

- The input data (vector  $x$ ) is arranged in a  $\sqrt{N} \times \sqrt{N}$  matrix (in row-major order);
- The matrix is distributed in **1D block-cyclic (column)** manner to the processors;

- The algorithm works in three steps:

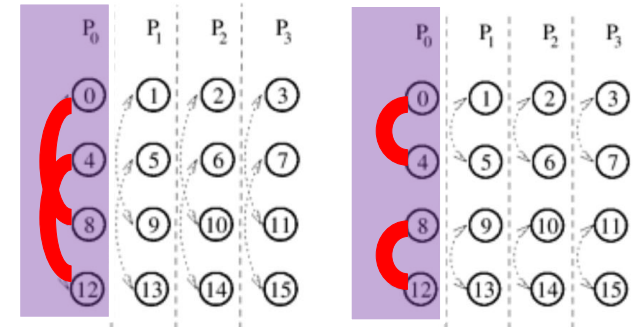
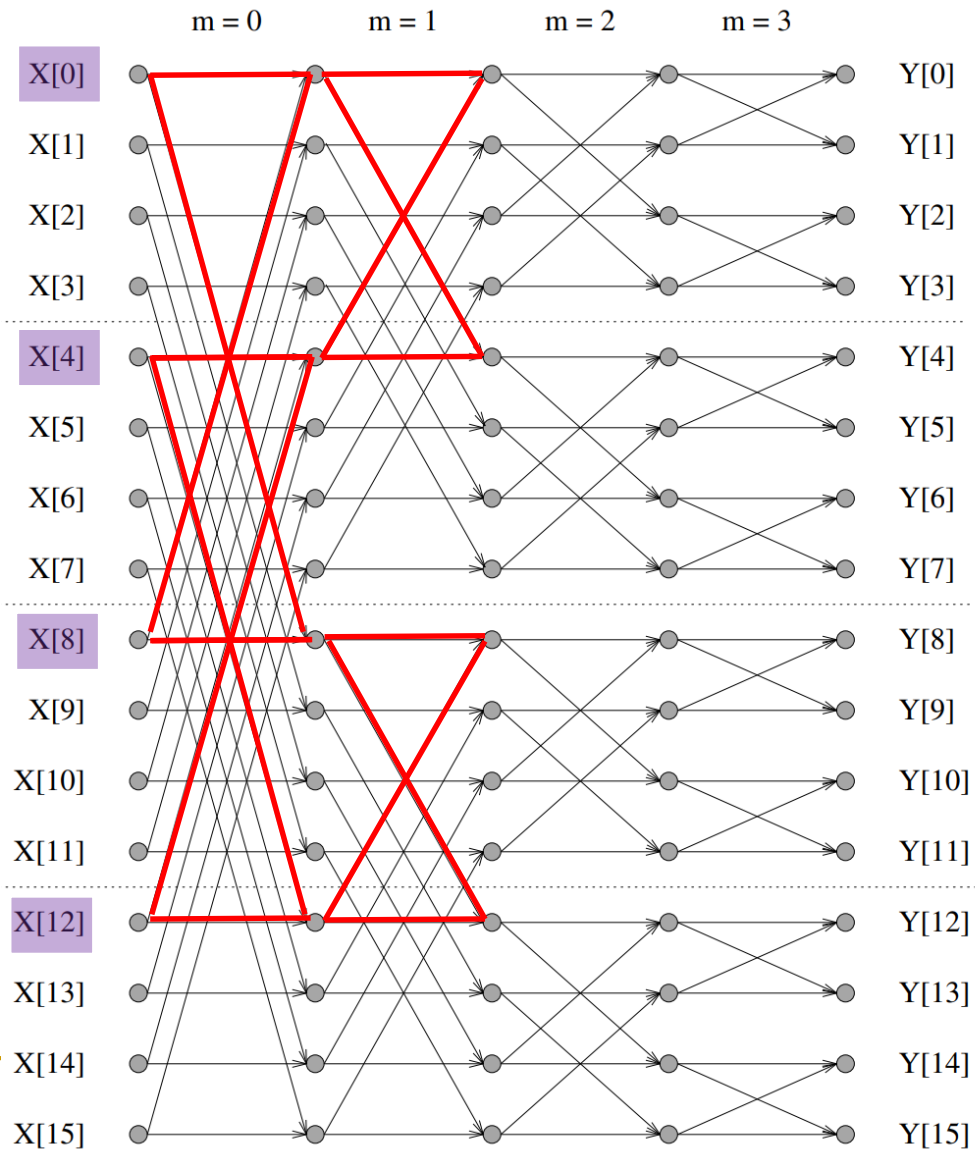
- 1) Each processor does  $\sqrt{N}/P$  local FFTs for each column of the matrix it holds;
- 2) Perform a global matrix transpose through an all-to-all personalized communication (MPI\_Alltoall);
- 3) Each processor again does  $\sqrt{N}/P$  local FFTs for each column of the new matrix it holds;

## (2) Transpose Algorithm



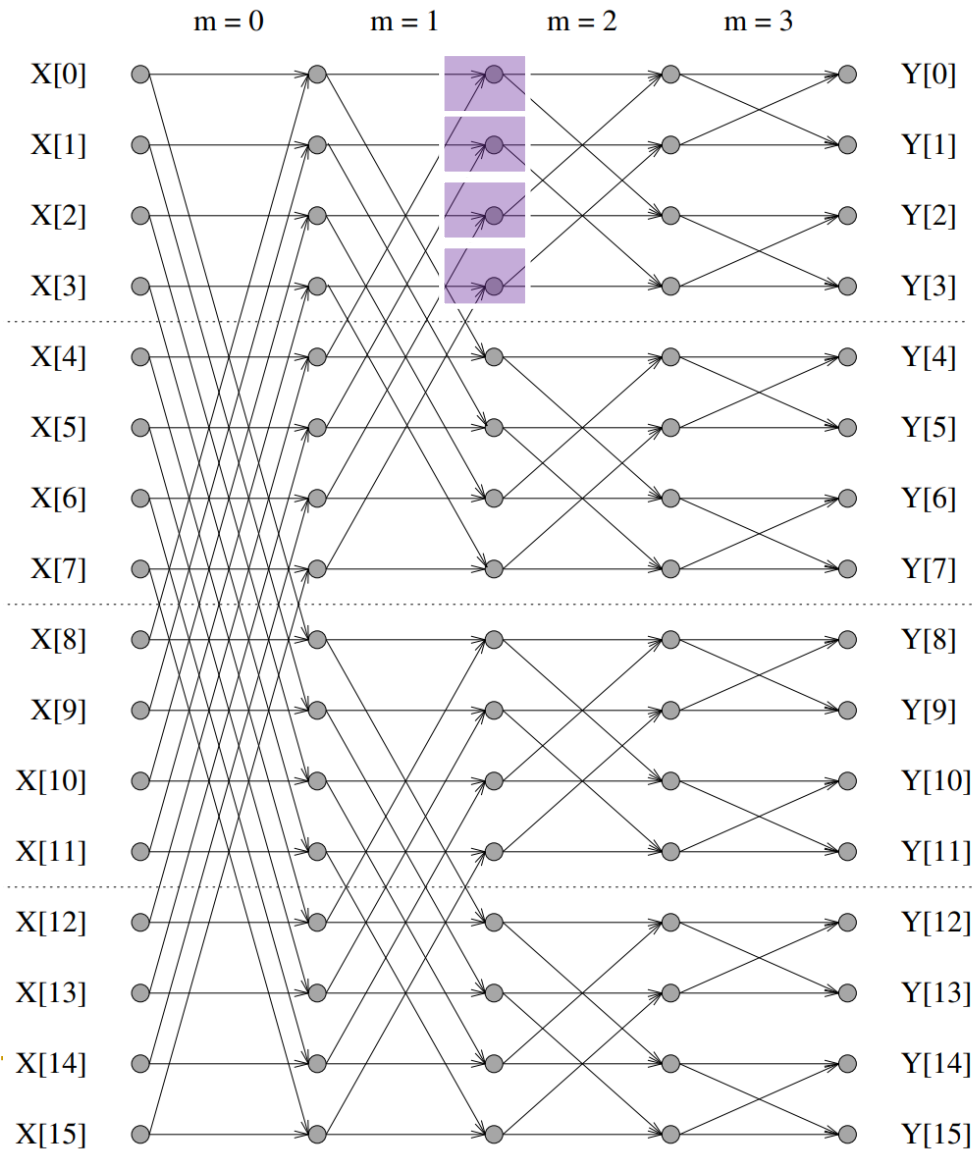
$P_0$	$P_1$	$P_2$	$P_3$
0	1	2	3
4	5	6	7
8	9	10	11
12	13	14	15

## (2) Transpose Algorithm



**Step 1:** Local FFT for each column (on same processor)

## (2) Transpose Algorithm



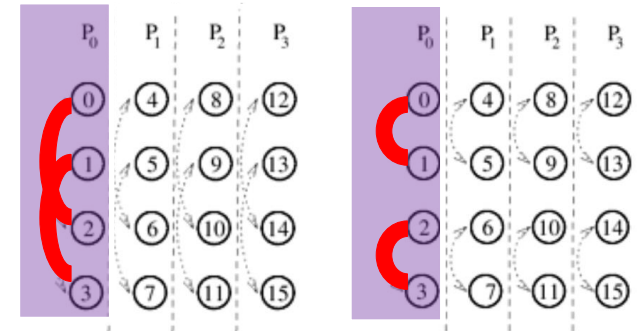
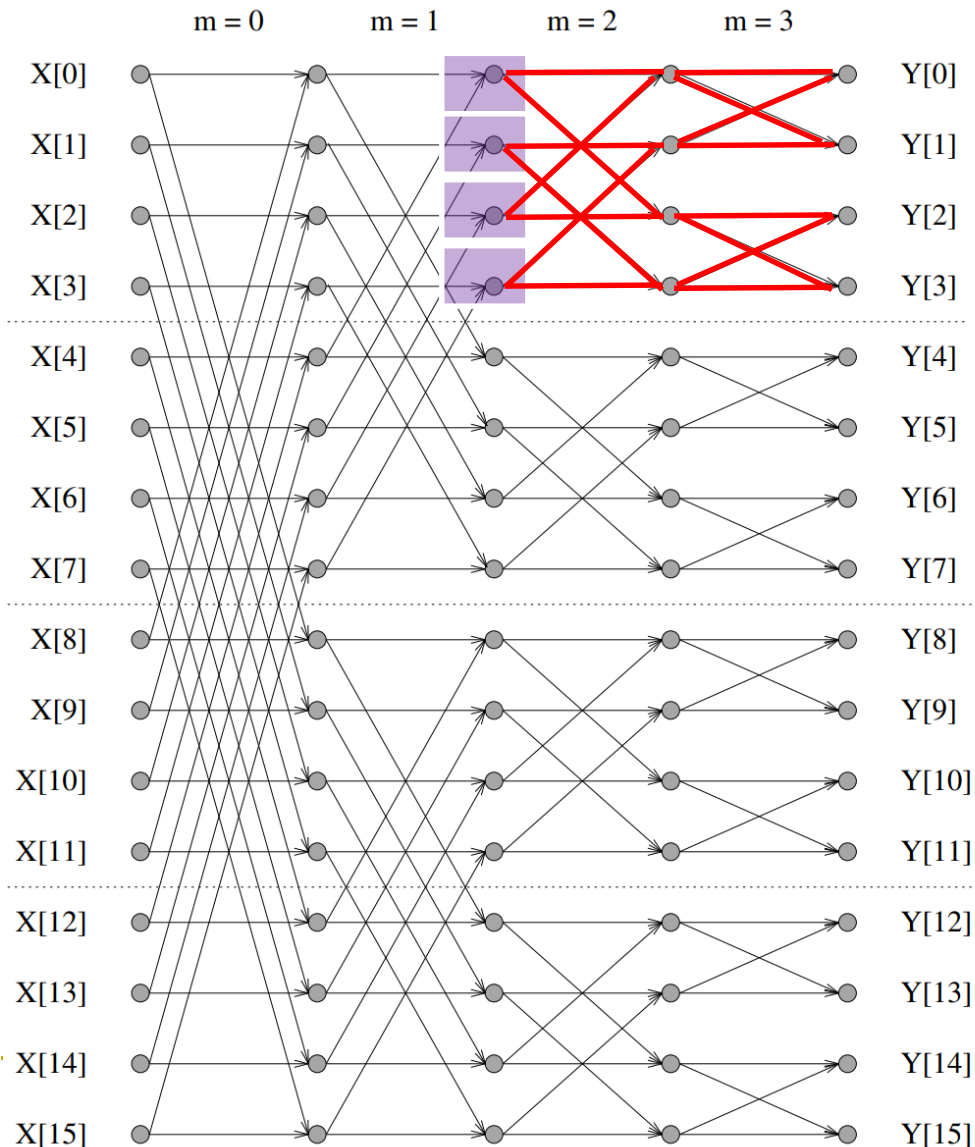
$P_0$	$P_1$	$P_2$	$P_3$
0	4	8	12
1	5	9	13
2	6	10	14
3	7	11	15

**Step 1:** Local FFT for each column (on same processor)

**Step 2:** Transpose matrix (all-to-all personalized comm)



## (2) Transpose Algorithm



**Step 1:** Local FFT for each column (on same processor)

**Step 2:** Transpose matrix (all-to-all personalized comm)


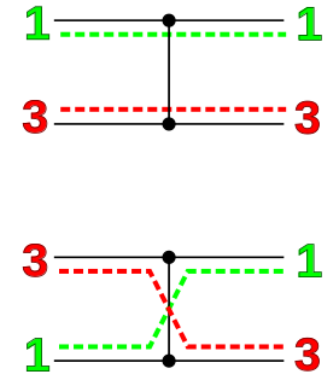
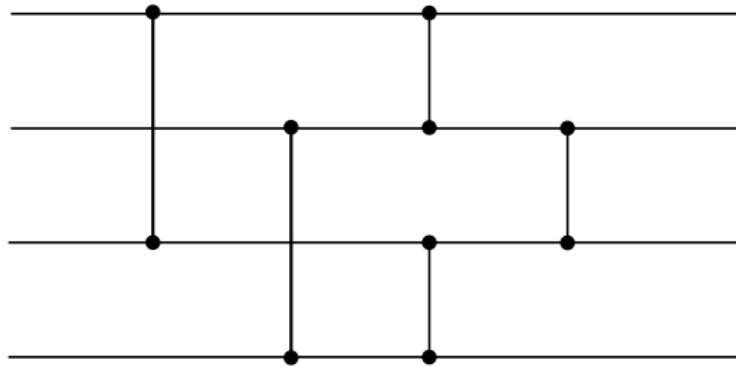
**Step 3:** Local FFT for each column (on same processor)

# Sorting Algorithms

- A **sorting algorithm** arranges a list of  $n$  numbers or items in ascending or descending order --- a very important procedure in scientific computing.
- We have seen previously **quicksort** and **mergesort** in shared-memory architectures.
- In this class, we discuss three other sorting algorithms and their parallel implementations:
  - 1) **Bitonic Sort;**
  - 2) **Brick Sort;**
  - 3) **Shear Sort.**

# Sorting Networks

- A **sorting network** sorts a fixed number of input values using a network of **comparators** connected by **wires**.

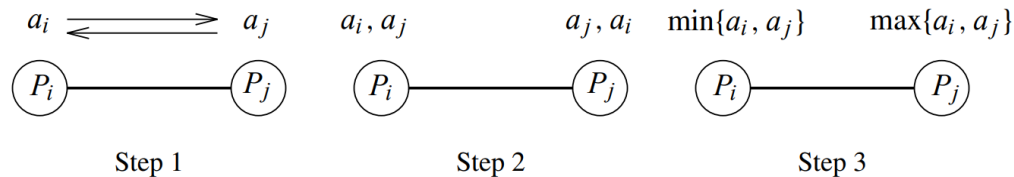


***Why is this sorting network able to sort any four input values in ascending order?***

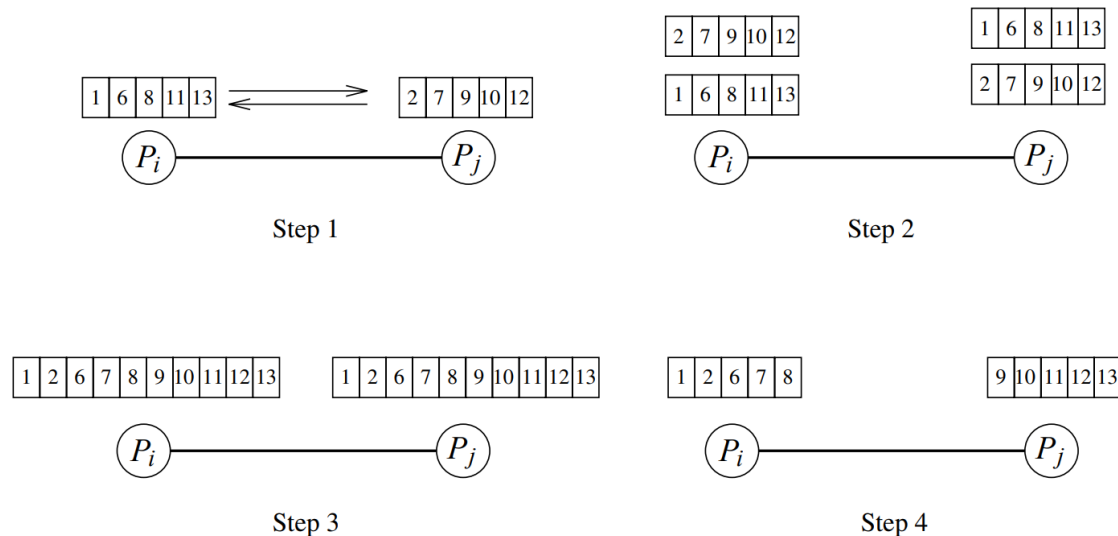
Figure source: [https://en.wikipedia.org/wiki/Sorting\\_network](https://en.wikipedia.org/wiki/Sorting_network)

# Basic Comparison Operations in Parallel

## ■ Compare-Exchange (one element per process)



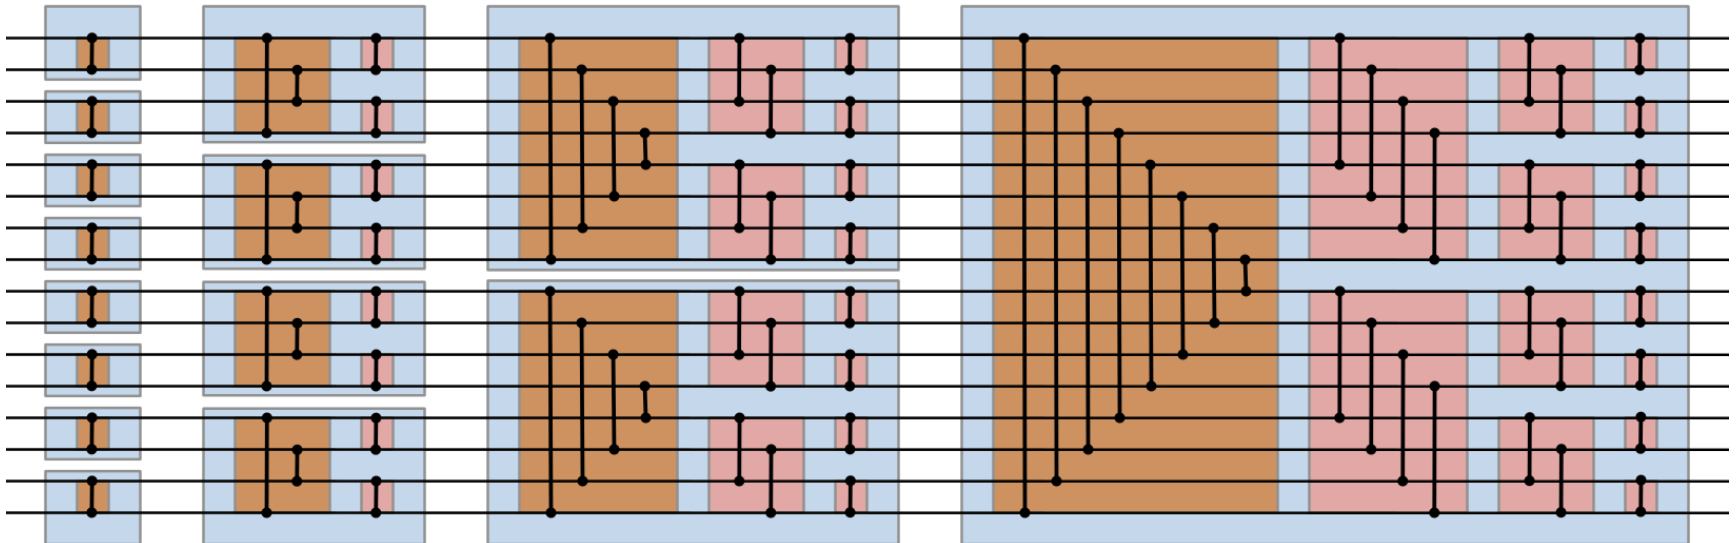
## ■ Compare-Split (multiple elements per process)



# (1) Bitonic Sort

## ■ Bitonic Sorting Network

- Constructed recursively using **divide-and-conquer** for any number of inputs  $n$  that is a power of 2 (*read the reference materials on details of its construction*).



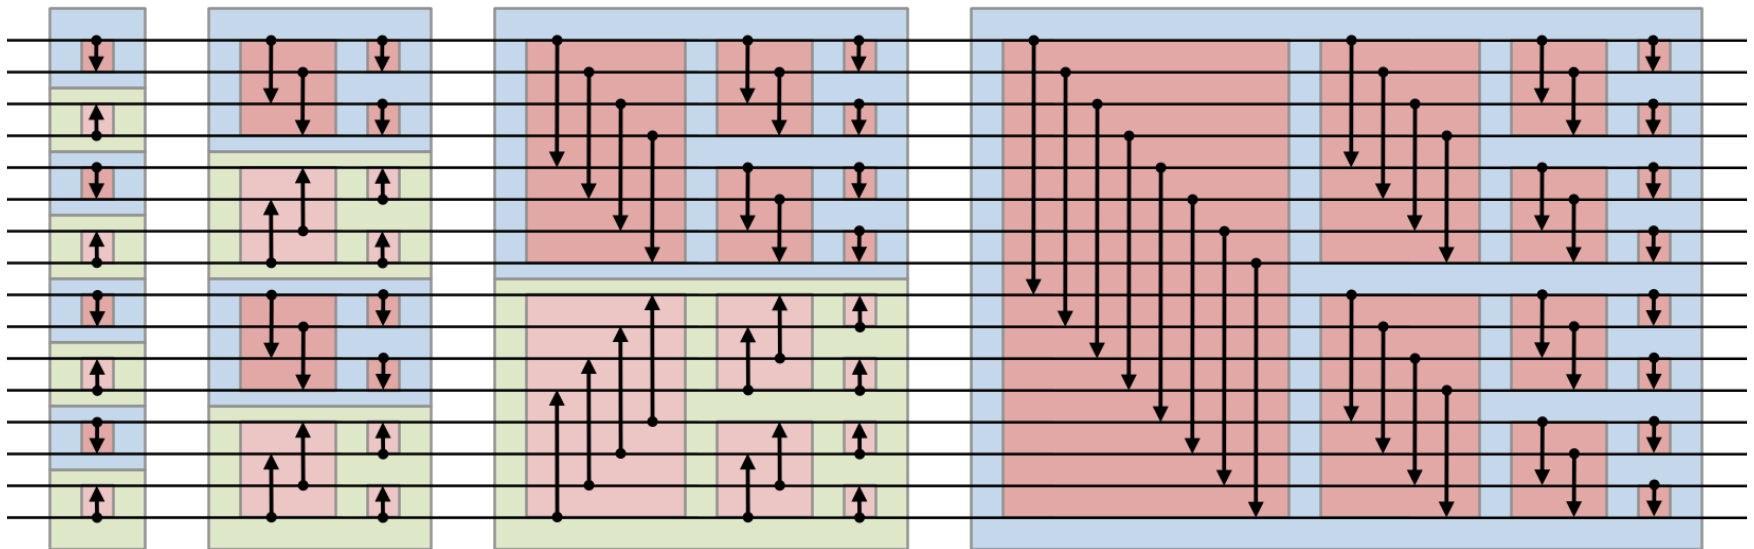
**Standard form**: *the larger number of a comparator always goes down.*

Figure source: [https://en.wikipedia.org/wiki/Bitonic\\_sorter](https://en.wikipedia.org/wiki/Bitonic_sorter)

# (1) Bitonic Sort

## ■ Bitonic Sorting Network

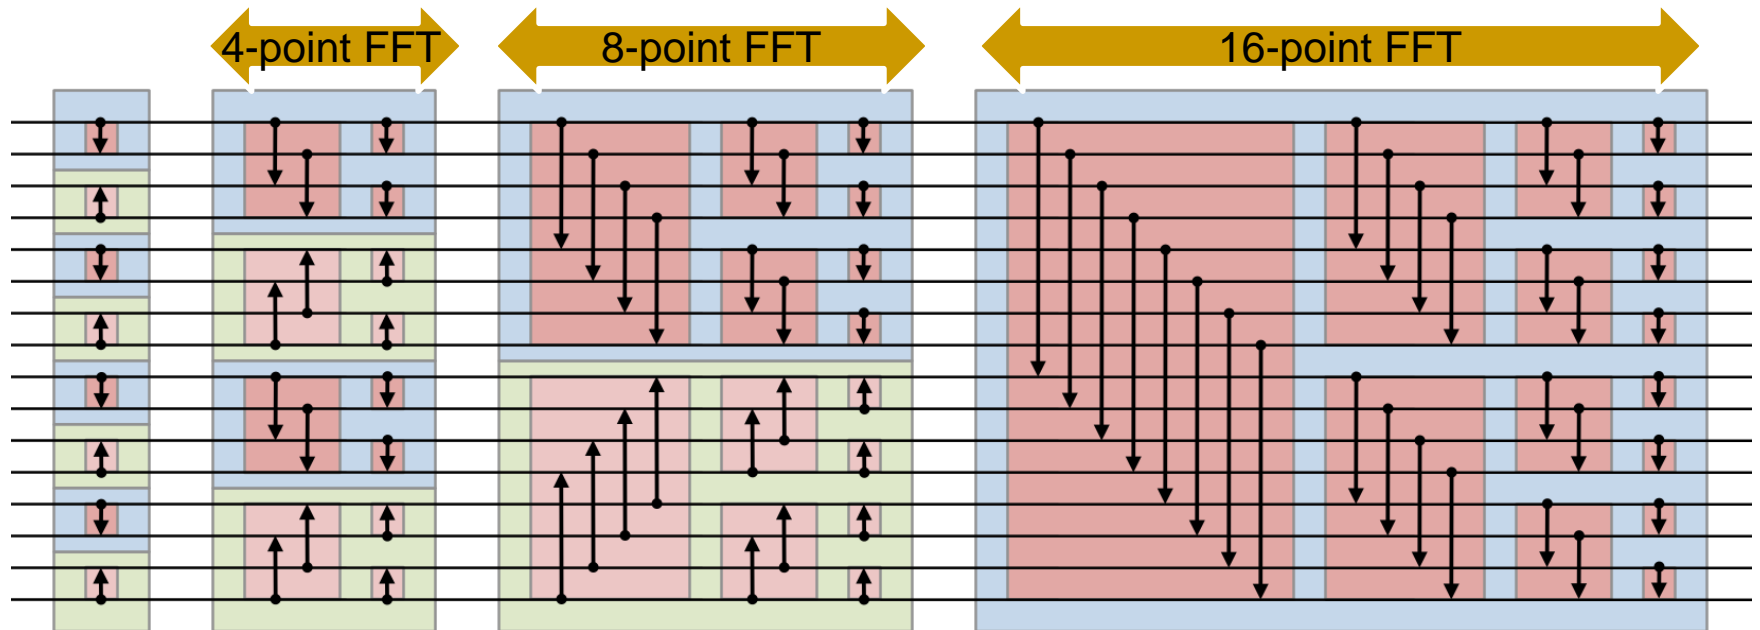
- Constructed recursively using **divide-and-conquer** for any number of inputs  $n$  that is a power of 2 (*read the reference materials on details of its construction*).



**Non-standard form**: *the larger number of a comparator goes to point of arrow.*

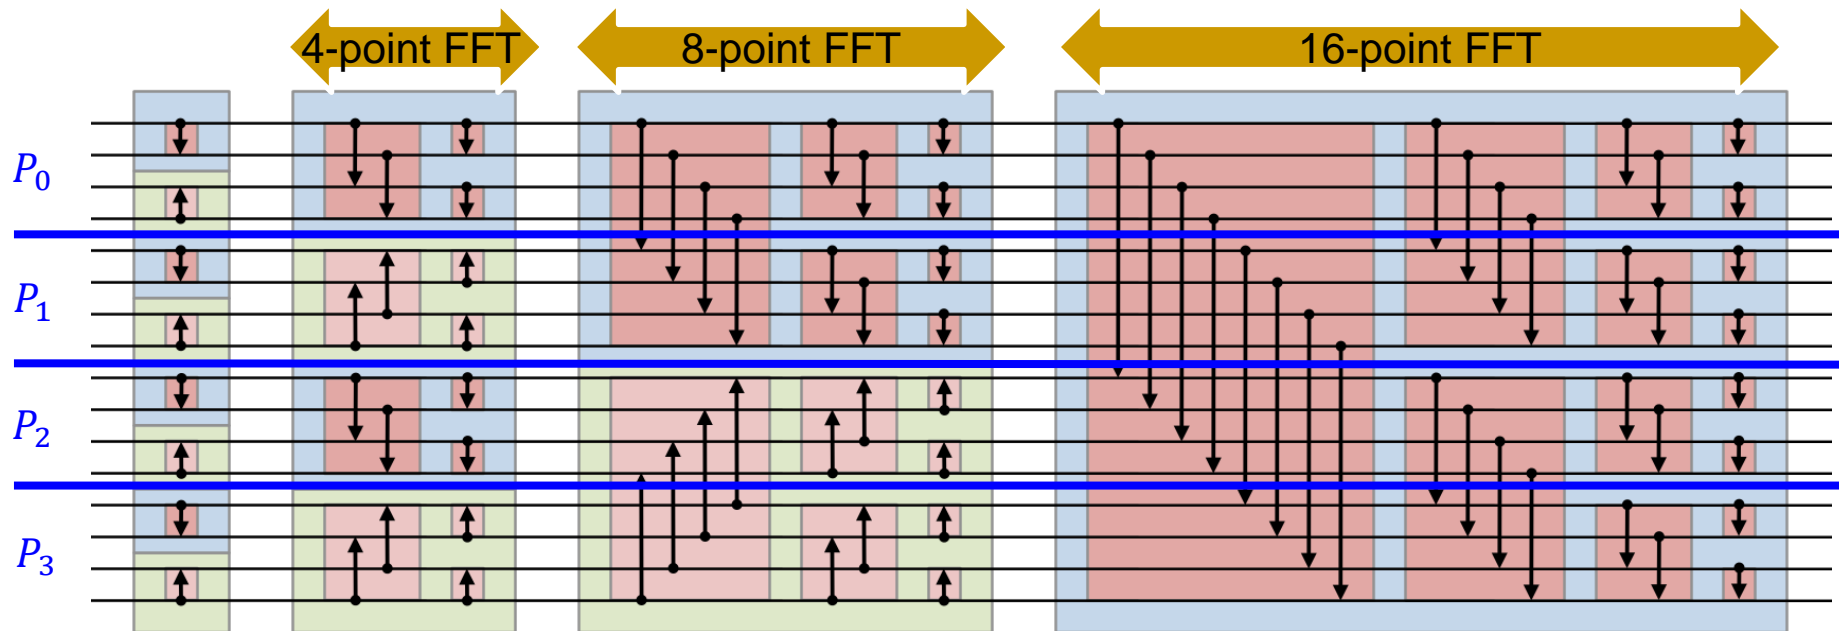
# (1) Bitonic Sort in Parallel

Observe that a bitonic sorting network (non-standard) has multiple stages of computations with **FFT-like structures**.



# (1) Bitonic Sort in Parallel

Observe that a bitonic sorting network (non-standard) has multiple stages of computations with **FFT-like structures**.



Parallel algorithms similar to Binary-Exchange for FFT can be implemented for Bitonic Sort.

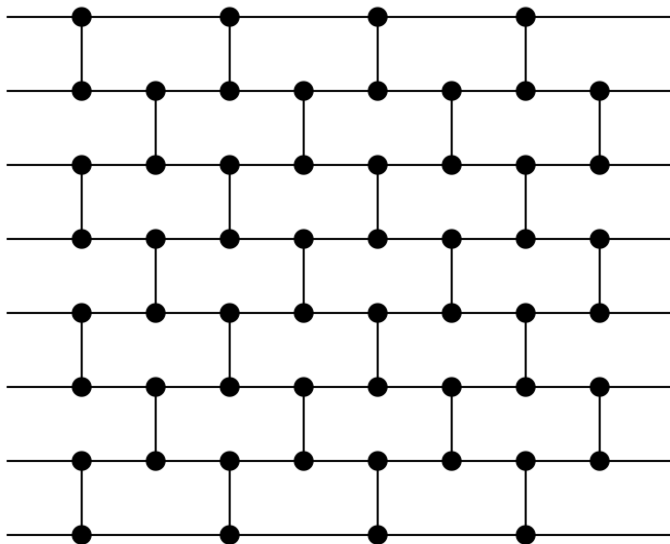
- *1D Block distribution* of data, and comparisons across two processors can be done via *compare-exchange* operations.



## (2) Brick Sort

### ■ Odd-Even Sorting Network

- Constructed by  $n$  levels of comparators connected in a “brick-like” pattern, with alternating odd and even sorting steps (it’s also called *Odd-Even Transposition Sort*).



Unsorted array: 2, 1, 4, 9, 5, 3, 6, 10

Step 1(odd): 2 1 4 9 5 3 6 10

Step 2(even): 1 2 4 9 3 5 6 10

Step 3(odd): 1 2 4 3 9 5 6 10

Step 4(even): 1 2 3 4 5 9 6 10

Step 5(odd): 1 2 3 4 5 6 9 10

Step 6(even): 1 2 3 4 5 6 9 10

Step 7(odd): 1 2 3 4 5 6 9 10

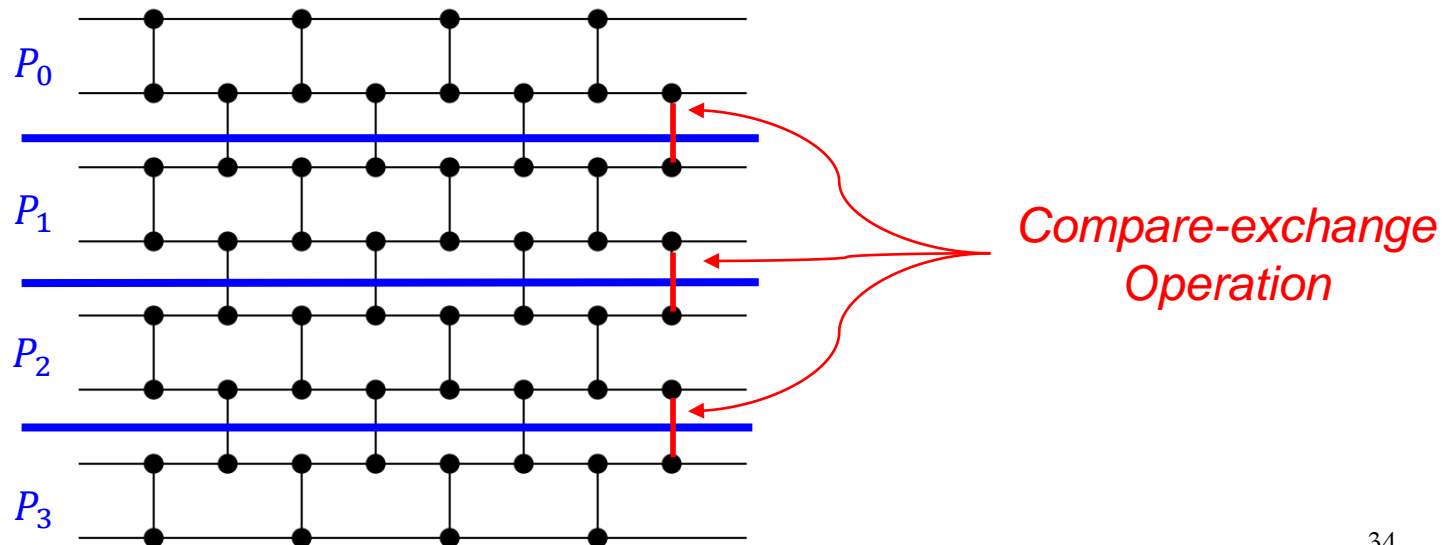
Step 8(even): 1 2 3 4 5 6 9 10

Sorted array: 1, 2, 3, 4, 5, 6, 9, 10

## (2) Brick Sort in Parallel

### ■ Parallel Implementation 1:


- ❑ 1D Block distribution;
- ❑ Follow exactly the same pattern of an  $n$ -step odd-even sorting network;
- ❑ Only boundary elements of a processor need communication (done via *compare-exchange*).



## (2) Brick Sort in Parallel

### ■ Parallel Implementation 2:

- 1D Block distribution;
- Follow the same pattern again, but treating blocks of elements together  $\rightarrow P$  steps instead of  $n$ ;
- Entire block of elements of a processor need to be communicated and compared (done via *compare-split*).



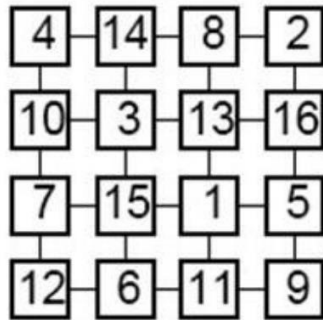
### (3) Shear Sort

- Algorithm works in 2D:
  - A list of  $n$  numbers initially arranged in a  $\sqrt{n} \times \sqrt{n}$  matrix;
  - For step  $k = 1, 2, \dots, \lg n + 1$ :
    - If  $k$  is odd, sort odd rows in *ascending* order and sort even rows in *descending* order;
    - If  $k$  is even, sort all columns in *ascending* order;
  - The final list will be sorted in **2D snake-like** order.

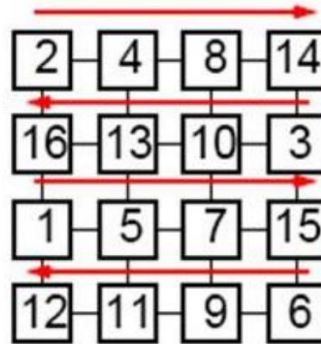
4	14	8	2
10	3	13	16
7	15	1	5
12	6	11	9

Original placement  
of numbers

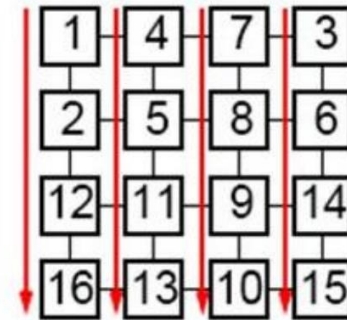
# Shear Sort



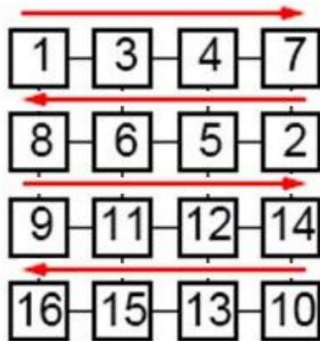
(a) Original placement of numbers




(b) Phase 1 — Row sort



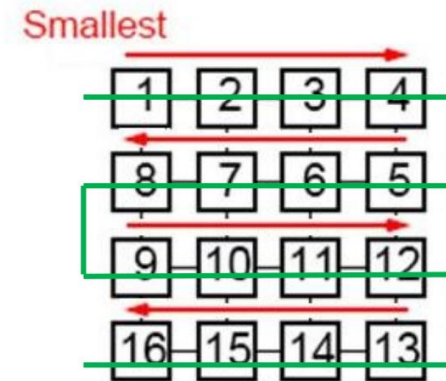
(c) Phase 2 — Column sort



(d) Phase 3 — Row sort



(e) Phase 4 — Column sort

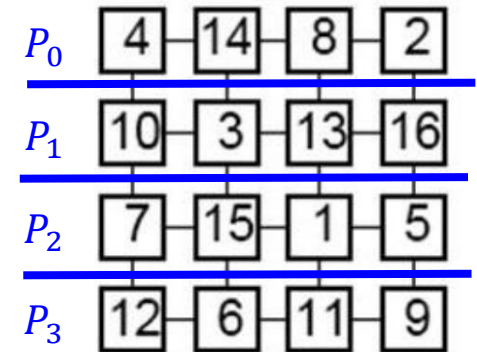


(f) Final phase — Row sort

# Shear Sort in Parallel

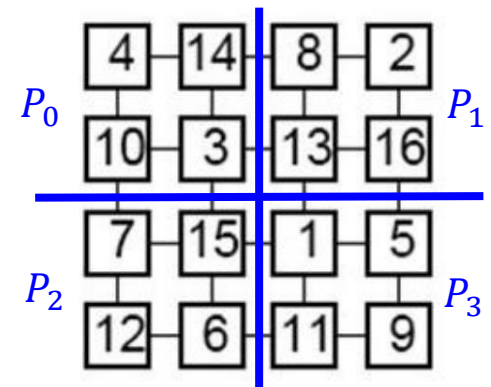
## ■ 1D Block Implementation:

- Row-sorting can be done *locally* on each processor;
- Column-sorting can be done across processors with *Brick Sort* for each column.



## ■ 2D Block Implementation:

- Both row-sorting and column-sorting can be done with *Brick Sort* across multiple processor groups (rows and columns).



# References

- A. Grama, A. Gupta, G. Karypis, V. Kumar. Introduction to Parallel Computing (2<sup>nd</sup> Edition). *Addison-Wesley Professional*.
- H. Casanova, A. Legrand, Y. Robert. Parallel Algorithms. *Chapman & Hall/CRC*.
- B. Barney. Introduction to Parallel Computing Tutorial, Lawrence Livermore National Laboratory.  
[https://computing.llnl.gov/tutorials/parallel\\_comp/](https://computing.llnl.gov/tutorials/parallel_comp/)